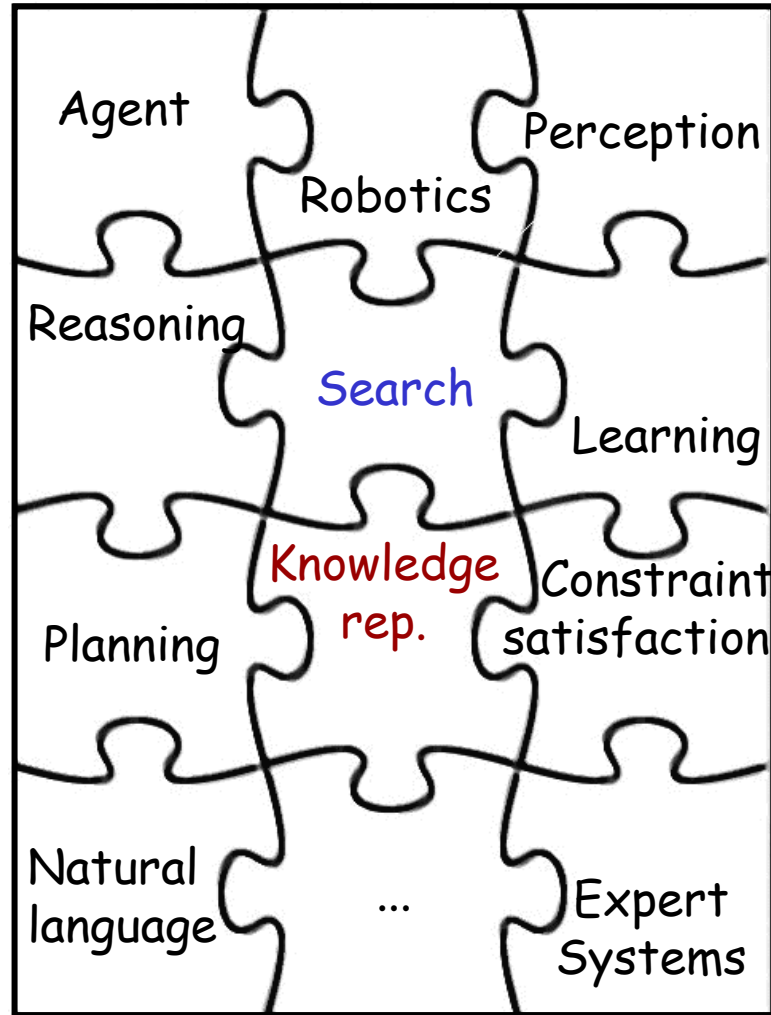
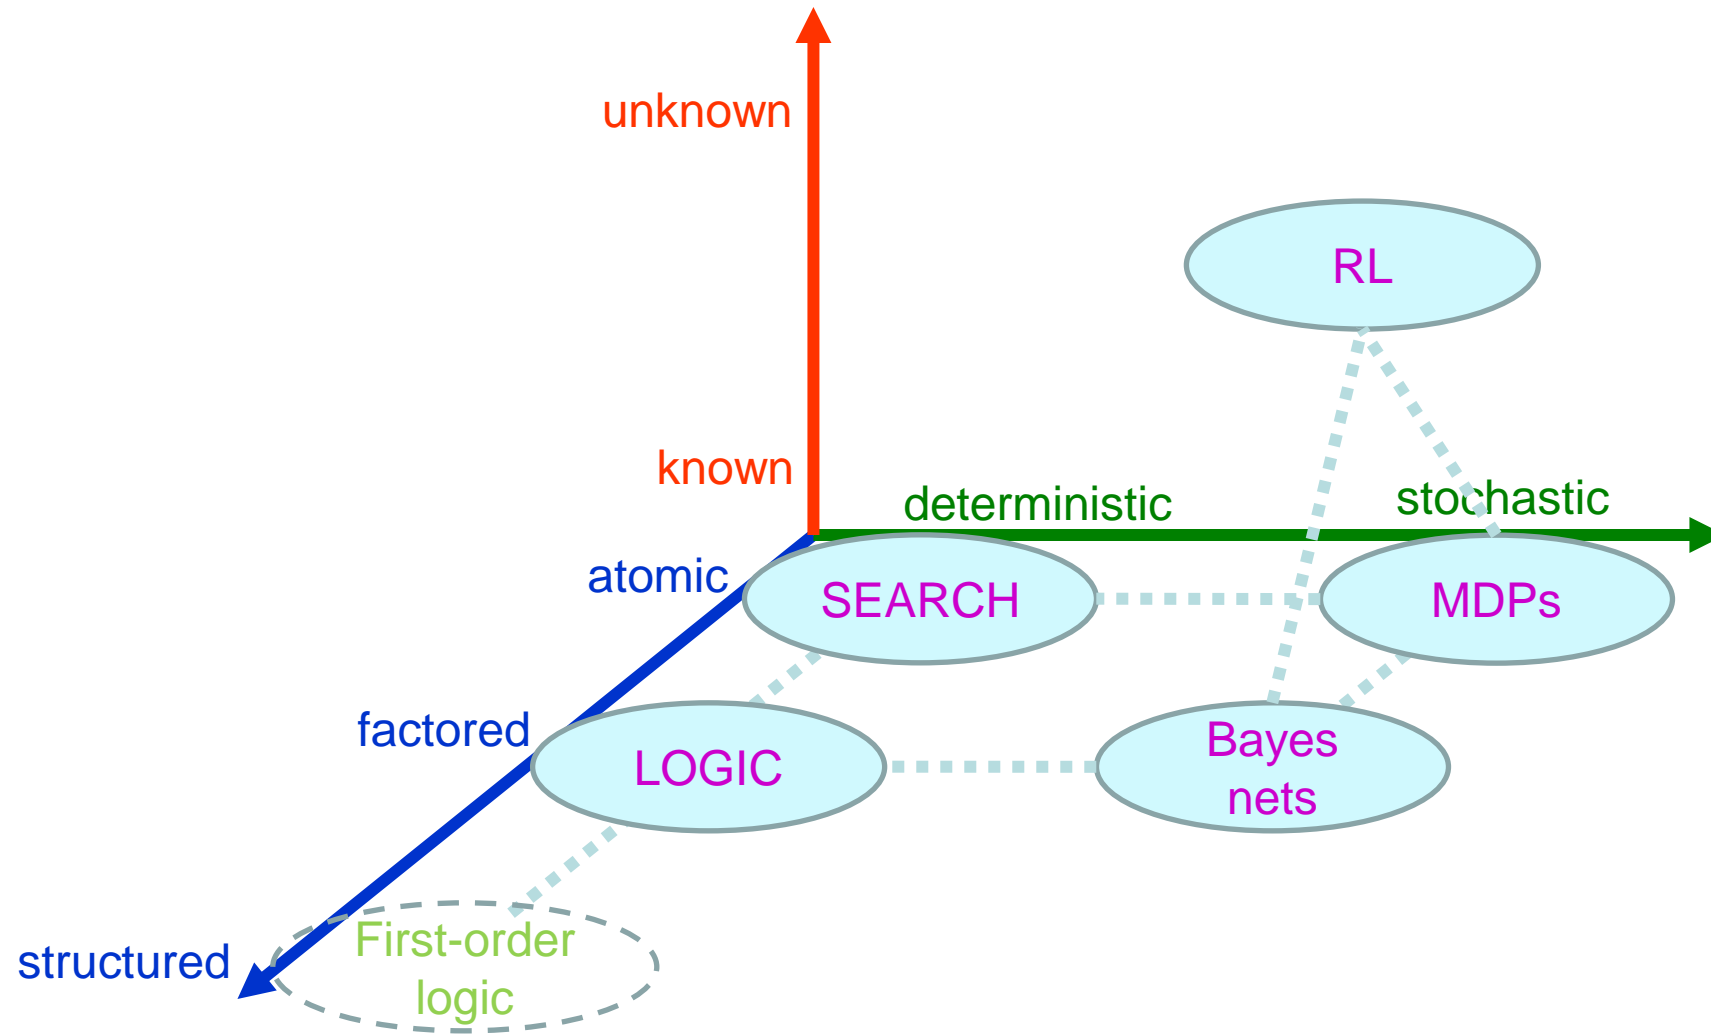


Artificial Intelligence




Outline of AI Strategies




Why Game Examples

- Across history, puzzles and games requiring the exploration of alternatives have been considered a challenge for human intelligence:
- Chess originated in Persia and India about 4000 years ago
- Checkers appear in 3600-year-old Egyptian paintings
- Go originated in China over 3000 years ago
- So, it's not surprising that AI uses games to design and test algorithms



Artificial Intelligence

Search




Motivation

- One of the major goals of AI is to help humans in solving complex tasks
 - How can I fill my container with pallets?
 - Which is the *shortest* way from Pune to Delhi?
 - Which is the *fastest* way from Pune to Delhi?
 - How can I *optimize* the load of my cargo to *maximize* my revenue?
 - How can I solve my Sudoku game?
 - What is the sequence of actions I should apply to win a game?
- Sometimes **finding a solution is not enough**, you want the **optimal solution** according to some “cost” criteria
- All the example presented involve looking for a **plan**
- **Plan:** the **set of operations** to be performed of an **initial state**, to reach a final state that is considered the **goal state**
- Thus, we **need efficient techniques to search for paths**, or sequences of actions, that can enable us to reach the goal state, i.e., to find a plan
- Such techniques are commonly called ***Search Methods***


Today

- Agents that Plan Ahead
- Search Problems
- Uninformed Search Methods
 - Depth-First Search
 - Breadth-First Search
 - Uniform-Cost Search



Planning Agents

- Planning agents decide based on evaluating future action sequences
- Must have a model of how the world evolves in response to actions
- Usually have a definite goal
- Optimal: Achieve goal at least cost



Problem solving

(using intelligence similar to human intelligence or using Intelligent Agents)

- We want:

- To automatically solve a problem.

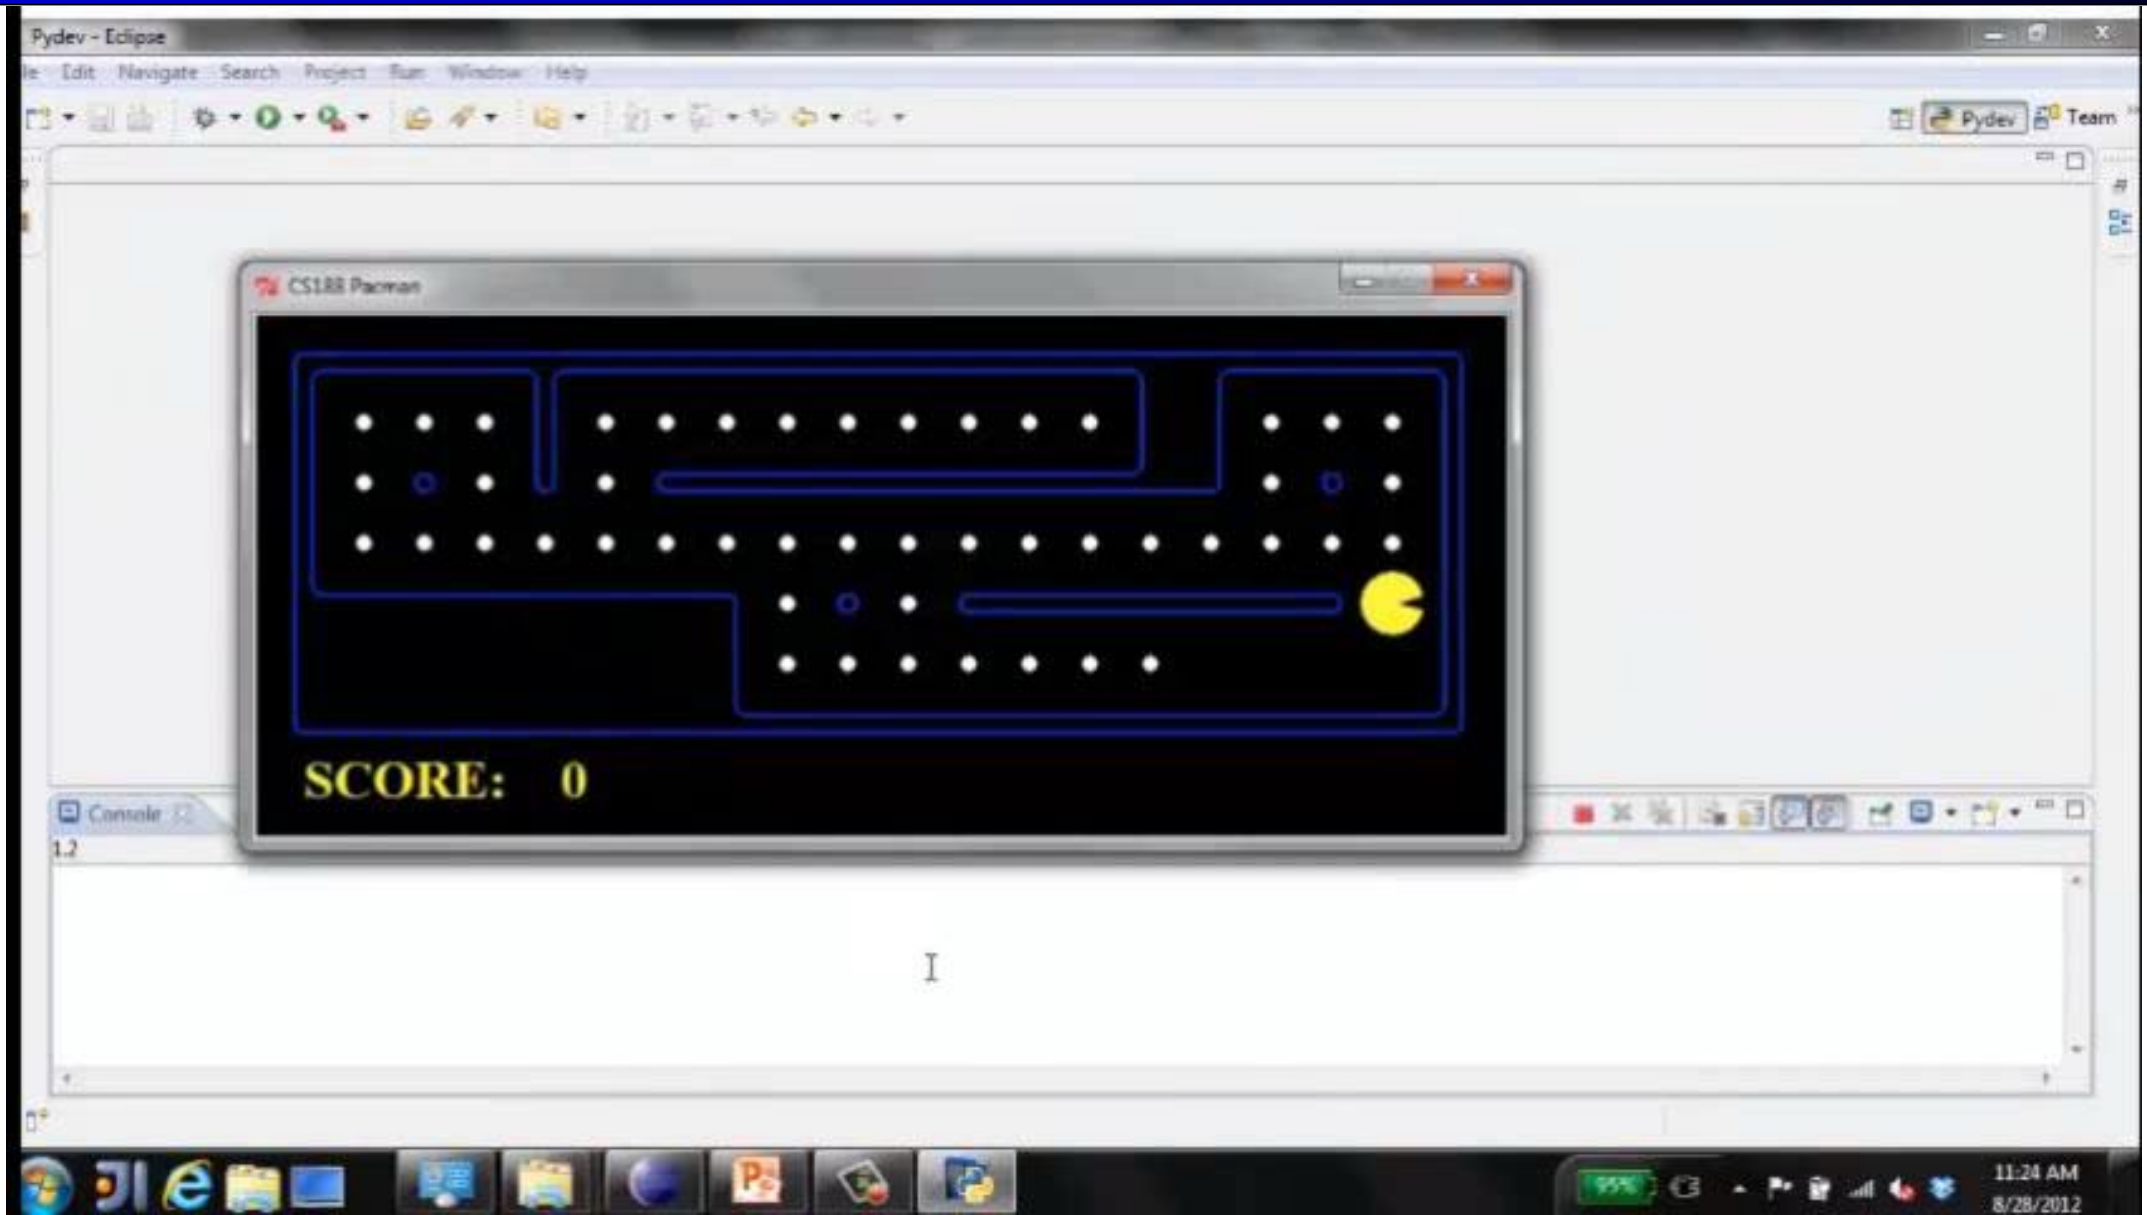
- We need:

- A representation of the problem.
- Algorithms that use some strategy to solve the problem defined in that representation.


Formalization

Searching technique

Move to nearest dot and eat it



Precompute optimal plan, execute it




Search Problems



Search Problems


- A search problem consists of:

- A state space \mathcal{S}
- An initial state s_0
- Actions $\mathcal{A}(s)$ in each state
- Transition model $Result(s, a)$
- A goal test $G(s)$
 - S has no dots left
- Action cost $c(s, a, s')$
 - -1 per step; +10 food; +500 win; -500 die; +200 eat ghost



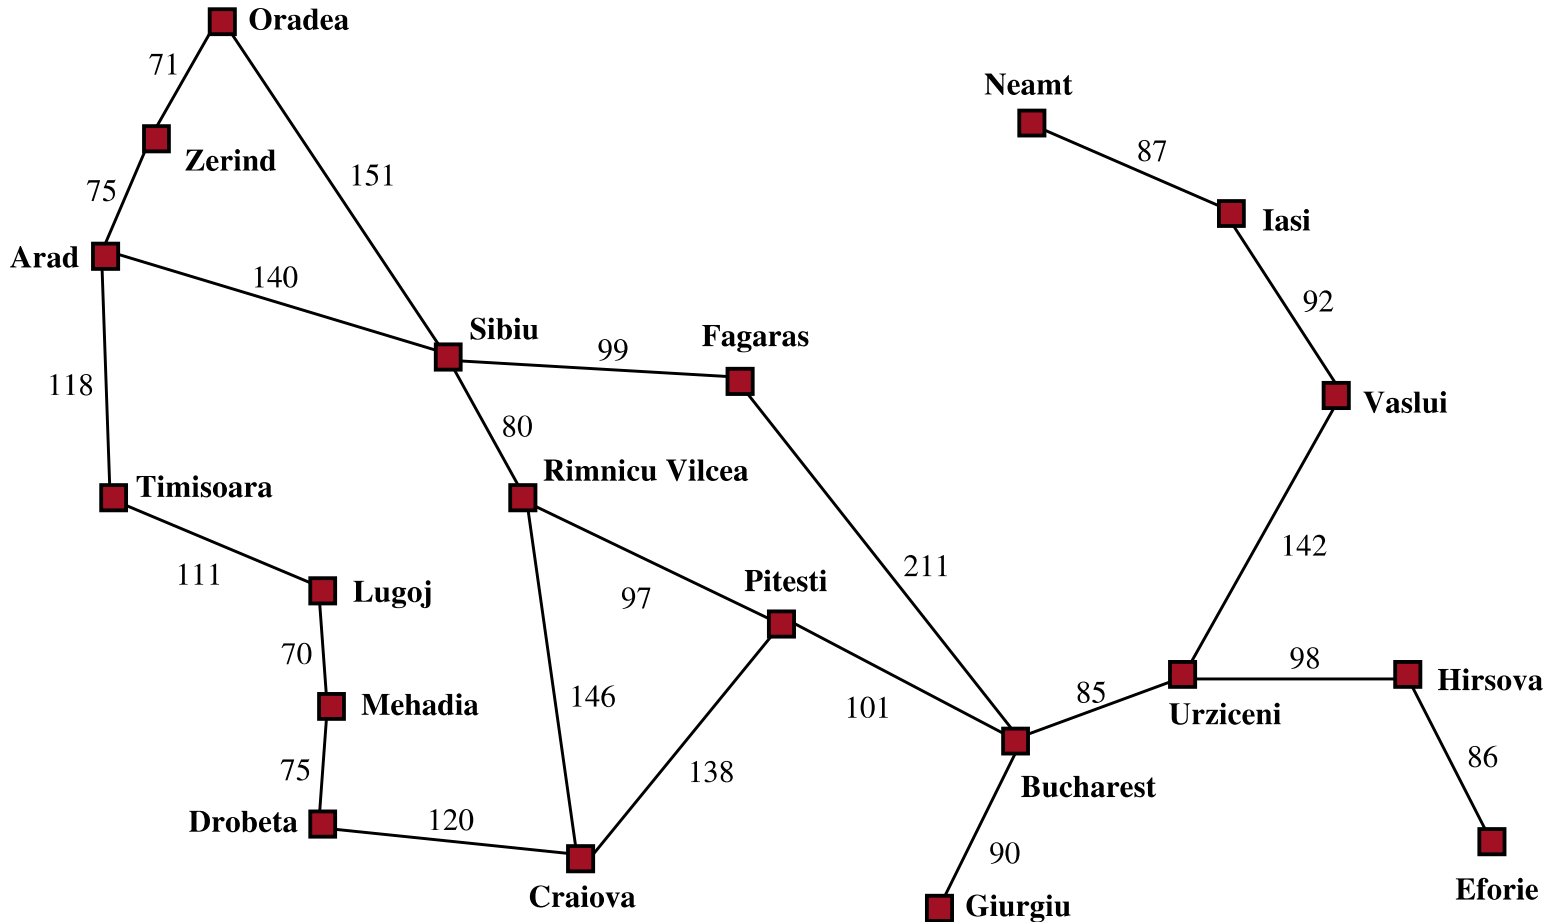
- A solution is an action sequence that reaches a goal state
- An optimal solution has least cost among all solutions

Example-1: Vacuum Cleaner world




- **States:** Discrete: dirt and robot location
- **Initial state:** Any
- **Actions:** *Left, right, suck*
- **Goal test:** No dirt at all locations
- **Path cost:** 1 per action
- **Solution** Optimal sequence of operations(actions)

Example-2: Traveling in Romania



- State space:
 - Cities
- Initial state:
 - Arad
- Actions:
 - Go to adjacent city
- Transition model:
 - Reach adjacent city
- Goal test:
 - $s = \text{Bucharest?}$
- Action cost:
 - Road distance from s to s'
- Solution?

Example-3: 8-Puzzle



Search is about the
exploration of alternatives

Example-3: 8-Puzzle

8	2	
3	4	7
5	1	6

Initial state

1	2	3
4	5	6
7	8	

Goal state

State: Any arrangement of 8 numbered tiles and an empty tile on a 3x3 board


- **States:** Locations of **tiles**
- **Initial State:** Given
- **Actions:** Move blank left, right, up, down
- **Goal test:** Goal state (given)
- **Path cost:** 1 per move
- **Solution:** Optimal sequence of operators

How big is the state space of the (n^2-1) -puzzle?

- 8-puzzle \rightarrow ?? states

How big is the state space of the (n^2-1) -puzzle?

- 8-puzzle \rightarrow ?? states




How big is the state space of the (n^2-1) -puzzle?

- 8-puzzle $\rightarrow 9! = 362,880$ states
- 15-puzzle $\rightarrow 16! \sim 2.09 \times 10^{13}$ states
- 24-puzzle $\rightarrow 25! \sim 10^{25}$ states

But only half of these states are reachable from any given state
(but you may not know that in advance)

8-, 15-, 24-Puzzles




How big is the state space?


Tic-Tac-Toe



Connect Four



Ultimate Tic-Tac-Toe



Chess



<9 possible moves/turn
255,168 possible games


~7 possible moves/turn
~4.5 x 10¹² possible games

~9 possible moves/turn
~4.4 x 10³⁸ possible games

~37 possible moves/turn
~10¹²⁰ possible games


Search Space Representation

- A graph is *undirected* if arcs do not imply a direction, *direct* otherwise
- A graph is *connected* if every pair of nodes is connected by a path
- A connected graph with no loop is called *tree*
- A *weighted graph*, is a graph for which a value is associated to each arc




State Graph

- Each state is represented by a distinct node
- An arc (or edge) connects a node s to a node s' if $s' \in \text{SUCCESSORS}(s)$
- The state graph may contain more than one connected component




Solution to the Search Problem

- A **solution** is a path connecting the initial node to a goal node (any one)




Solution to the Search Problem

- A **solution** is a path connecting the initial node to a goal node (any one)
- The **cost** of a path is the sum of the arc costs along this path
- An **optimal** solution is a solution path of minimum cost
- There might be no solution !




Searching the State Space





- It is often not feasible (or too expensive) to build a complete representation of the state graph

Searching the State Space





- Often it is not feasible (or too expensive) to build a complete representation of the state graph
- A problem solver must construct a solution by exploring a small portion of the graph



Searching the State Space





Searching the State Space





Searching the State Space





Searching the State Space





Searching the State Space



Searching the State Space




Search Tree



Note that some states may be
visited multiple times

State Space Graphs vs. Search Trees


State Space Graph



Each NODE in the search tree is an entire PATH in the state space graph.


We construct the tree on demand – and we construct as little as possible.

Search Tree



Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:




How big is its search tree (from S)?




Quiz: State Space Graphs vs. Search Trees

Consider this 4-state graph:



How big is its search tree (from S)?



Important: Those who don't know history are doomed to repeat it!


Search Trees

Open List / Fringe: [

]

Closed List / Visited: [

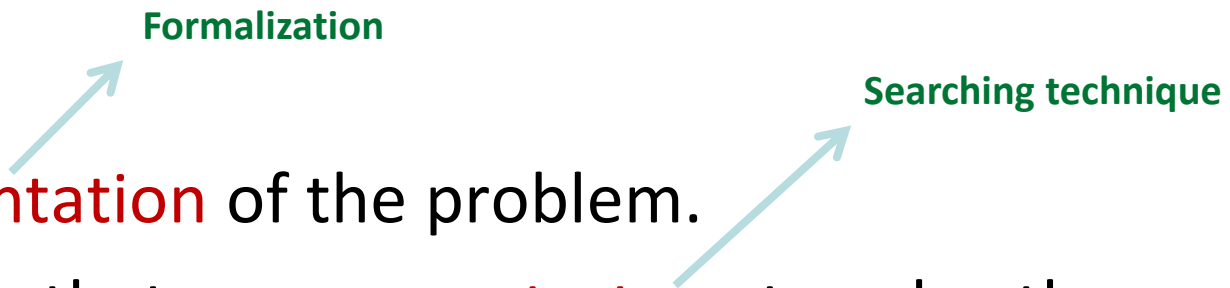
]



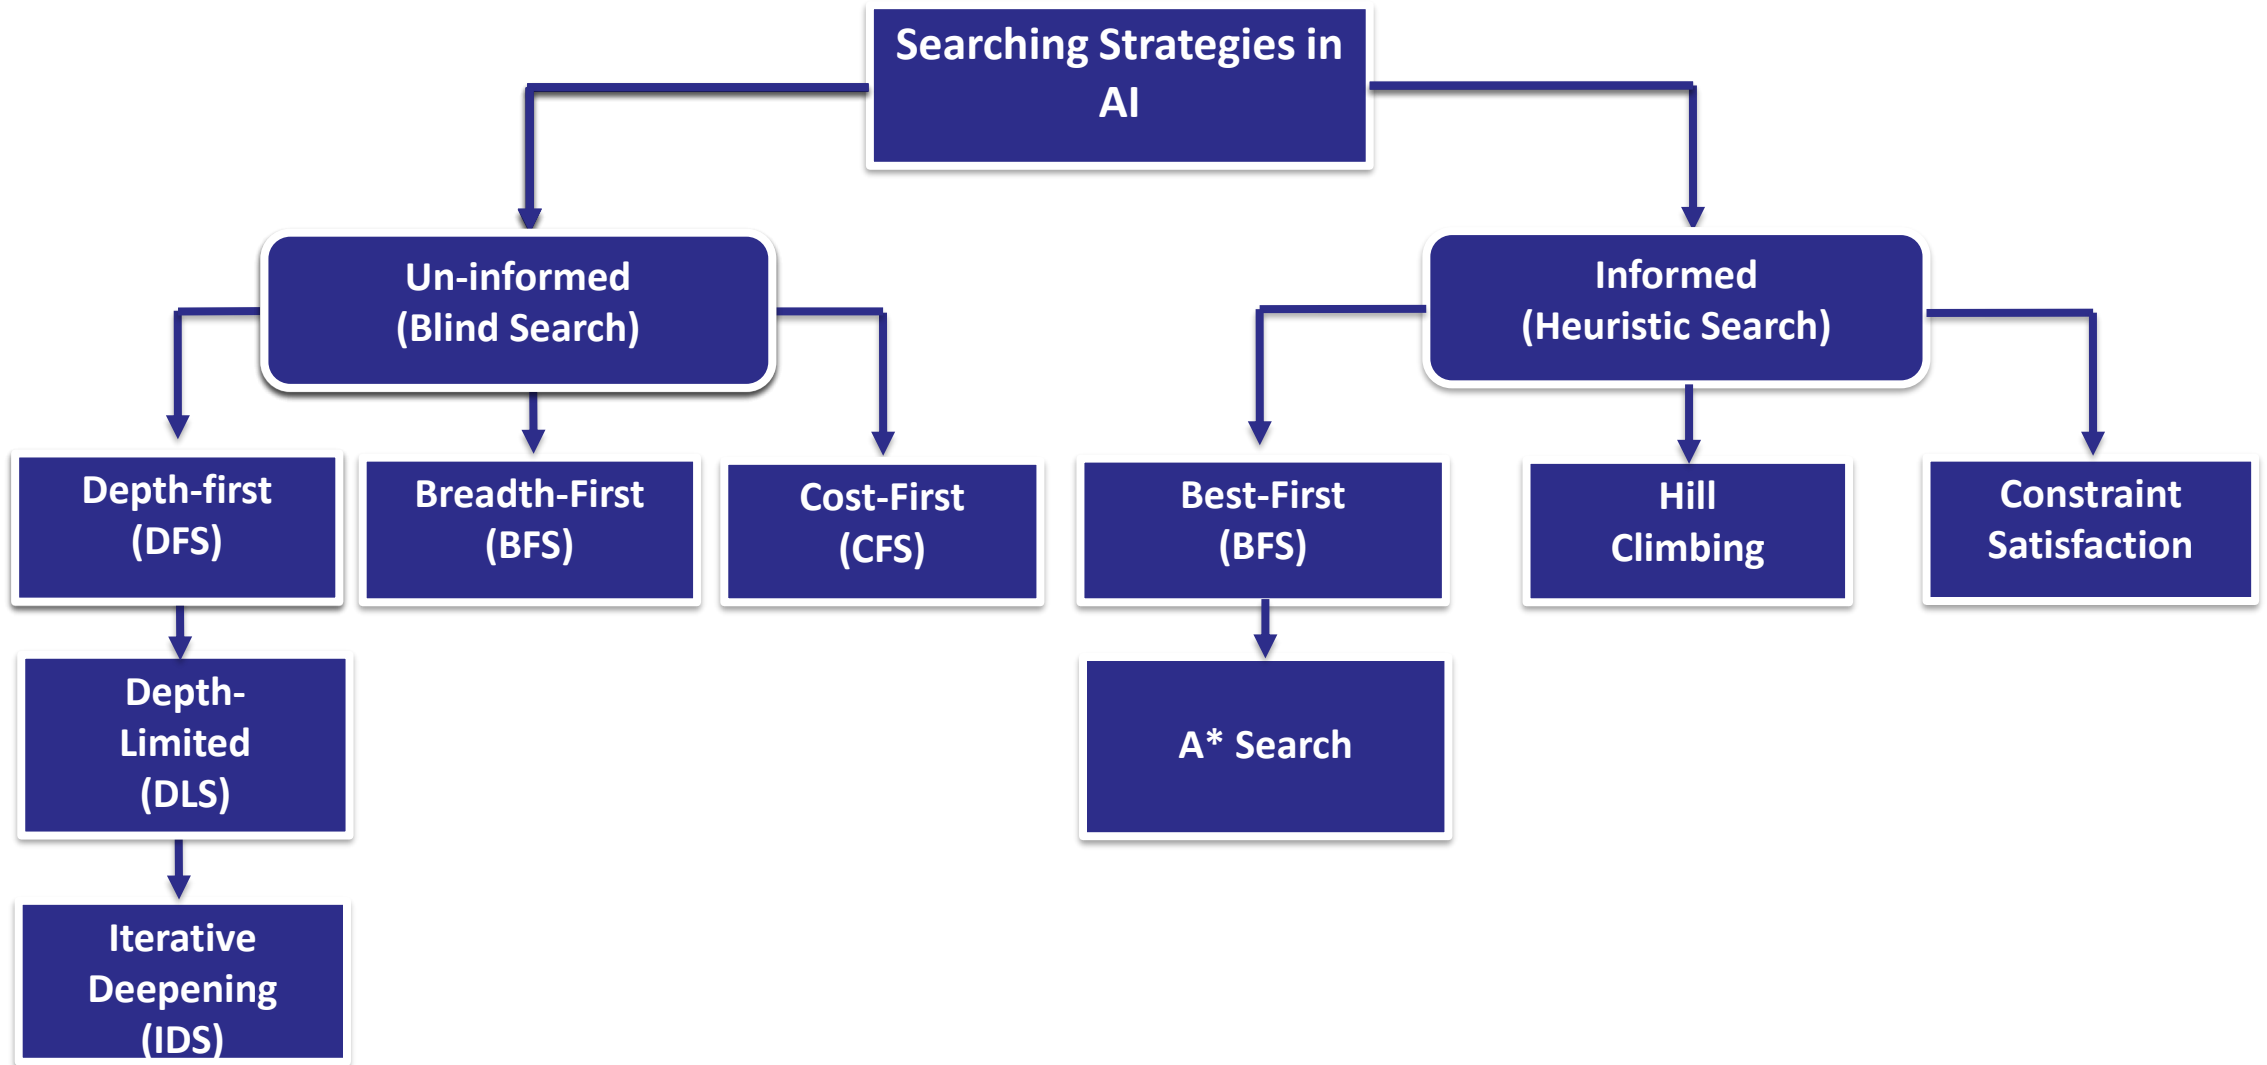
Search Problems Are Models



Problem solving

- We want:
 - To **automatically** solve a problem.
 - We need:
 - A **representation** of the problem.
 - Algorithms that use **some strategy** to solve the problem defined in that representation.
- 
- The diagram consists of two light blue arrows. The first arrow originates from the word 'representation' in the second bullet point and points diagonally upwards and to the right towards the word 'Formalization'. The second arrow originates from the words 'some strategy' in the third bullet point and points diagonally upwards and to the right towards the words 'Searching technique'.

AI searching Strategies






A very large number of AI problems are formulated as search problems.

Example

For a **blind strategy**, N_1 and N_2 are just two nodes
(at some position in the search tree)

Strategies do not exploit state descriptions to order
FRINGE






Goal state

Example


For a **heuristic strategy** counting the number of misplaced tiles, N_2 is more promising than N_1

Strategies exploit state descriptions to order FRINGE



Goal state

Data Structure of a Node




Depth of a node N = length of path from root to N
(depth of the root = 0)

Node expansion

The **expansion** of a node N of the search tree consists of:


- 1) Evaluating the successor function on $\text{STATE}(N)$
- 2) Generating a child of N for each state returned by the function


node generation \neq node expansion



Fringe of Search Tree

- The **fringe** is the set of all search nodes that haven't been expanded yet





Is it identical to
the set of
leaves?

Search Strategy

- Fringe:

Set of all search nodes that haven't been expanded yet

- Implemented as a priority queue FRINGE

- INSERT(node, FRINGE)
- REMOVE(FRINGE)

- The ordering of the nodes in FRINGE defines the search strategy

Performance Measures

- **Completeness:**

A search algorithm is complete if it finds a solution whenever one exists

[What about the case when no solution exists?]

- **Optimality:**

A search algorithm is optimal if it returns a minimum-cost path whenever a solution exists

- **Complexity:**

It measures the time and amount of memory required by the algorithm

General Graph Search

```
function GRAPH-SEARCH(problem, fringe) return a solution, or failure
  closed ← an empty set
  fringe ← INSERT(MAKE-NODE(INITIAL-STATE[problem]), fringe)
  loop do
    if fringe is empty then return failure
    node ← REMOVE-FRONT(fringe)
    if GOAL-TEST(problem, STATE[node]) then return node
    if STATE[node] is not in closed then
      add STATE[node] to closed
      for child-node in EXPAND(STATE[node], problem) do
        fringe ← INSERT(child-node, fringe)
      end
    end
  end
end
```

- Main variations:
 - Which leaf node to expand next
 - Whether to check for repeated states
 - Data structures for frontier, expanded nodes

General Graph Search (Python)

INPUT: Start = S, Goal = G, SearchSpace = *dictionary*

OUTPUT: Status (True/False), Path, Goal

GraphSearch(SearchSpace, Start, Goal) :

 OpenList = [S]

 ClosedList = []

while *len*(OpenList) :

 N = *remove_first*(OpenList)

if N *not in* ClosedList : *insert*(N, ClosedList)

if N == Goal : *return* True, N

else:

 childrens = SearchSpace [N]

for child *in* childrens :

if child *not in* ClosedList *and* child *not in* OpenList :

 OpenList = *append*(OpenList, child)

return False

Note: this version is not saving the path for simplicity

Blind Strategies

- Breadth-first

- Bidirectional

- Depth-first

- Depth-limited
- Iterative deepening

- Uniform-Cost

(variant of breadth-first)

Arc cost = 1

Arc cost
= $c(\text{action}) \geq \epsilon > 0$



Depth-First Search




Depth-First Search

Strategy: expand a deepest node first

Implementation:
Fringe is a LIFO stack



Search Algorithm Properties



Search Algorithm Properties


- Complete: Guaranteed to find a solution if one exists?
- Optimal: Guaranteed to find the least cost path?
- Time complexity?
- Space complexity?

- Sketch of search tree:

- b is the branching factor
- m is the maximum depth
- solutions at various depths


- Number of nodes in entire tree?

- $1 + b + b^2 + \dots + b^m = O(b^m)$



Depth-First Search (DFS) Properties

- What nodes does DFS expand?
 - Some left prefix of the tree down to depth m .
 - Could process the whole tree!
 - If m is finite, takes time $O(b^m)$
- How much space does the frontier take?
 - Only has siblings on path to root, so $O(bm)$
- Is it complete?
 - m could be infinite
 - preventing cycles may help (more later)
- Is it optimal?
 - No, it finds the “leftmost” solution, regardless of depth or cost




Depth-Limited Search

- Depth-first with **depth cutoff** k (depth at which nodes are not expanded)
- Three possible outcomes:
 - Solution
 - Failure (no solution)
 - **Cutoff (no solution within cutoff)**

Iterative Deepening

- Idea: get DFS's space advantage with BFS's time / shallow-solution advantages
 - Run a DFS with depth limit 1. If no solution...
 - Run a DFS with depth limit 2. If no solution...
 - Run a DFS with depth limit 3.
- Isn't that wastefully redundant?
 - Generally, most work happens in the lowest level searched, so not so bad!



Iterative Deepening Search

Provides the best of both breadth-first and depth-first search

Main idea:

Totally horrifying !


IDS

For $k = 0, 1, 2, \dots$ do:


 Perform depth-first search with
 depth cutoff k

 (i.e., only generate nodes with depth $\leq k$)


Iterative Deepening



Iterative Deepening



Iterative Deepening



Performance

- Iterative deepening search is:
 - Complete
 - Optimal if step cost = 1
- Time complexity:
$$b^d + 2b^{d-1} + 3b^{d-2} + \dots + (d - 1)b^2 + db + (d + 1) = O(b^d)$$
- Space complexity: $O(bd)$ or $O(d)$

Number of Generated Nodes (Breadth-First & Iterative Deepening)

$d = 5$ and $b = 2$

d	BF	ID
0	1	$1 \times 6 = 6$
1	2	$2 \times 5 = 10$
2	4	$4 \times 4 = 16$
3	8	$8 \times 3 = 24$
4	16	$16 \times 2 = 32$
5	32	$32 \times 1 = 32$
SUM	63	120

$120/63 \sim 2$


Number of Generated Nodes (Breadth-First & Iterative Deepening)

$d = 5$ and $b = 10$

d	BF	ID
0	1	6
1	10	50
2	100	400
3	1,000	3,000
4	10,000	20,000
5	100,000	100,000
SUM	111,111	123,456

$123,456 / 111,111 \sim 1.111$



Breadth-First Search



Depth-First Search


Strategy: expand a shallowest node first

Implementation: Fringe is a FIFO queue




Breadth-First Search (BFS) Properties

- What nodes does BFS expand?
 - Processes all nodes above shallowest solution
 - Let depth of shallowest solution be s
 - Search takes time $O(b^s)$
- How much space does the frontier take?
 - Has roughly the last tier, so $O(b^s)$
- Is it complete?
 - s must be finite if a solution exists, so yes!
- Is it optimal?
 - If costs are equal (e.g., 1)




Bidirectional Strategy

2 fringe queues: FRINGE1 and FRINGE2




Time and space complexity is $O(b^{d/2}) \ll O(b^d)$
if both trees have the same branching factor b

Quiz: DFS vs BFS



Uniform Cost Search





Uniform Cost Search



$g(n)$ = cost from root to n

Strategy: expand lowest $g(n)$

Fringe is a priority queue sorted by $g(n)$




*Fringe is a priority queue
sorted by $g(n)$*



Uniform Cost Search (UCS) Properties

- What nodes does UCS expand?
 - Processes all nodes with cost less than cheapest solution!
 - If that solution costs C^* and arcs cost at least ϵ , then the “effective depth” is roughly C^*/ϵ
 - Takes time $O(b^{C^*/\epsilon})$ (exponential in effective depth)
- How much space does the frontier take?
 - Has roughly the last tier, so $O(b^{C^*/\epsilon})$
- Is it complete?
 - Assuming C^* is finite and $\epsilon > 0$, yes!
- Is it optimal?
 - Yes! (Proof next lecture via A*)



Comparison of Strategies

Measures	DFS	DLS	IDS	BFS	BDS	UCS
Time	b^d	b^l	b^d	b^d	$b^{\frac{d}{2}}$	b^d
Space	bm	b^l	bd	b^d	$b^{\frac{d}{2}}$	b^d
Optimal	No	No	Yes	Yes	Yes	Yes
Complete	No	$if\ l \geq d$	Yes	Yes	Yes	Yes