

# Informed (Heuristic) Search Strategies

- **Informed Search** – a strategy that uses problem-specific knowledge beyond the definition of the problem itself
- **Best-First Search** – an algorithm in which a node is selected for expansion based on an evaluation function  $f(n)$ 
  - Traditionally the node with the lowest evaluation function is selected
  - Not an accurate name...expanding the best node first would be a straight march to the goal.
  - Choose the node that *appears* to be the best

# Informed (Heuristic) Search Strategies

- There is a whole family of Best-First Search algorithms with different evaluation functions
  - Each has a heuristic function  $h(n)$
- $h(n)$  = estimated cost of the cheapest path from node  $n$  to a goal node
- Example: in route planning the estimate of the cost of the cheapest path might be the straight line distance between two cities

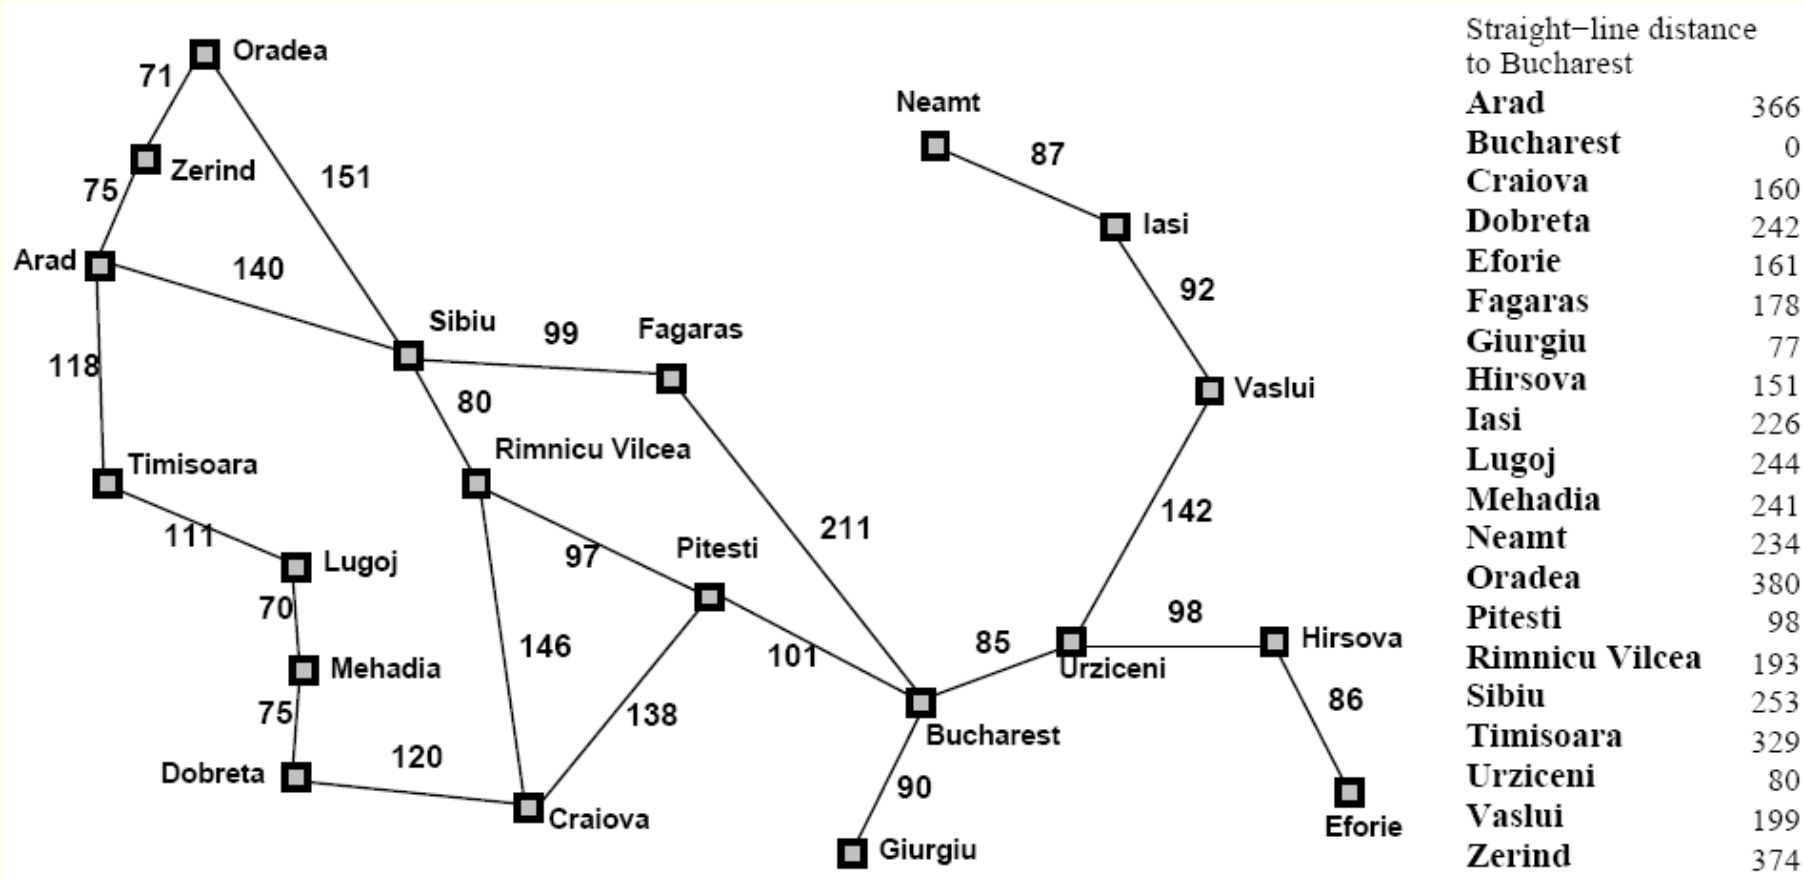
# A Quick Review

- $g(n)$  = cost from the initial state to the current state  $n$
- $h(n)$  = estimated cost of the cheapest path from node  $n$  to a goal node
- $f(n)$  = evaluation function to select a node for expansion (usually the lowest cost node)

# Greedy Best-First Search

- Greedy Best-First search tries to expand the node that is closest to the goal assuming it will lead to a solution quickly
  - $f(n) = h(n)$
  - aka “Greedy Search”
- Implementation
  - expand the “most desirable” node into the fringe queue
  - sort the queue in decreasing order of desirability
- Example: consider the straight-line distance heuristic  $h_{SLD}$ 
  - Expand the node that appears to be closest to the goal

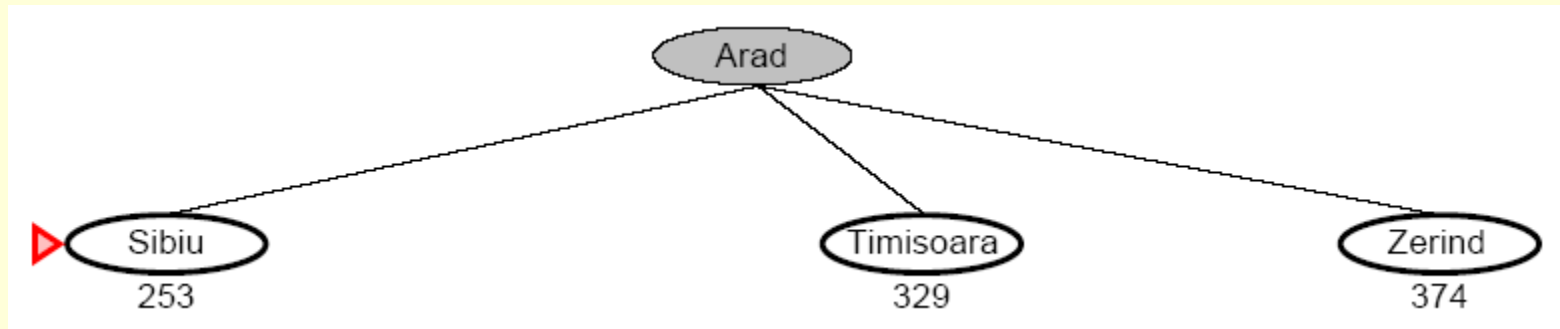
# Greedy Best-First Search



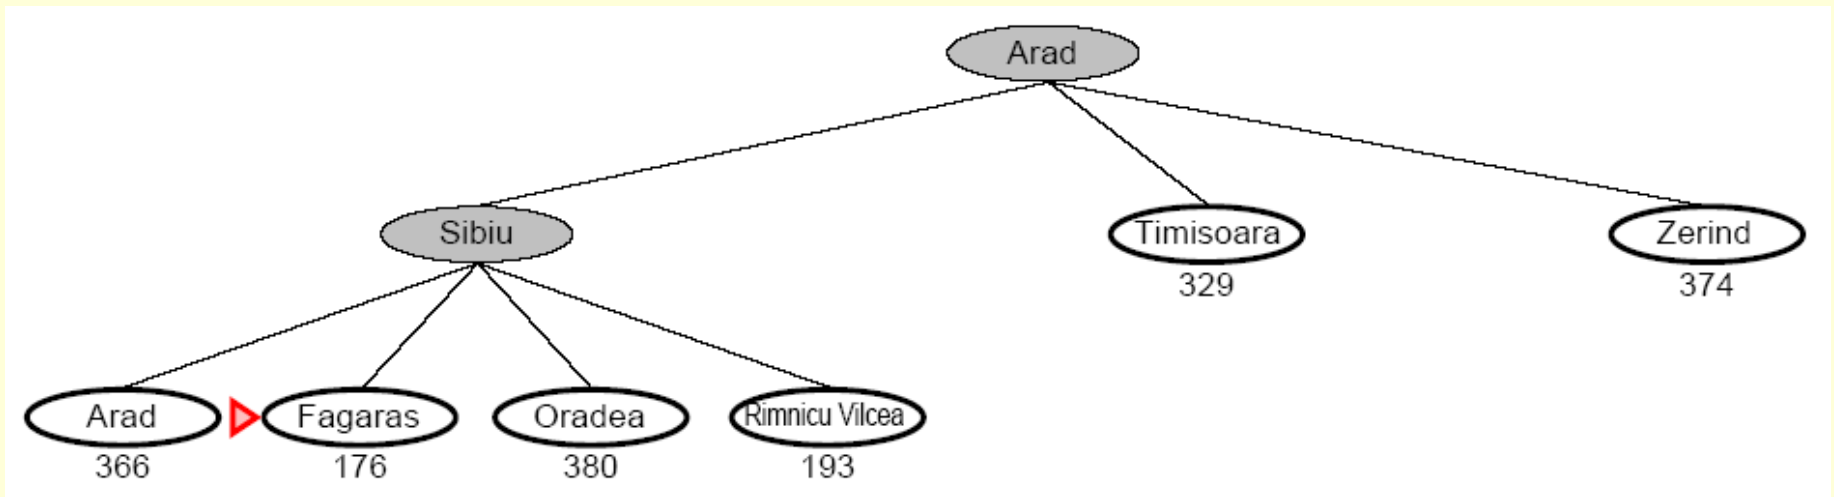
# Greedy Best-First Search

- $h_{\text{SLD}}(\text{In}(\text{Arid})) = 366$
- Notice that the values of  $h_{\text{SLD}}$  cannot be computed from the problem itself
- It takes some experience to know that  $h_{\text{SLD}}$  is correlated with actual road distances
  - Therefore a useful heuristic

# Greedy Best-First Search

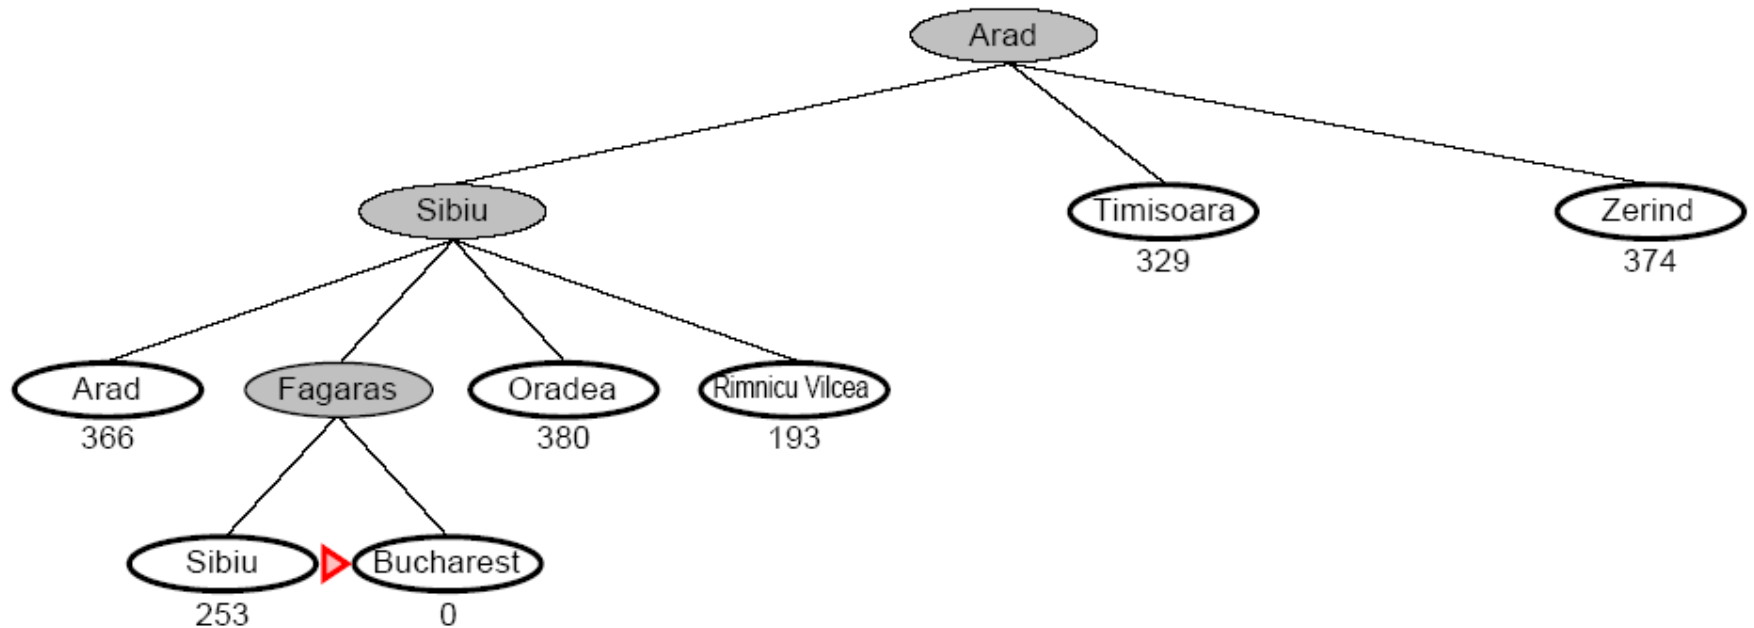


# Greedy Best-First Search





# Greedy Best-First Search



# Greedy Best-First Search

- Complete
  - No, GBFS can get stuck in loops (e.g. bouncing back and forth between cities)
- Time
  - $O(b^m)$  but a good heuristic can have dramatic improvement
- Space
  - $O(b^m)$  – keeps all the nodes in memory
- Optimal
  - No!

# A Quick Review - Again

- $g(n)$  = cost from the initial state to the current state  $n$
- $h(n)$  = estimated cost of the cheapest path from node  $n$  to a goal node
- $f(n)$  = evaluation function to select a node for expansion (usually the lowest cost node)

# A\* Search

- A\* (A star) is the most widely known form of Best-First search
  - It evaluates nodes by combining  $g(n)$  and  $h(n)$
  - $f(n) = g(n) + h(n)$
  - Where
    - $g(n)$  = cost so far to reach  $n$
    - $h(n)$  = estimated cost to goal from  $n$
    - $f(n)$  = estimated total cost of path through  $n$

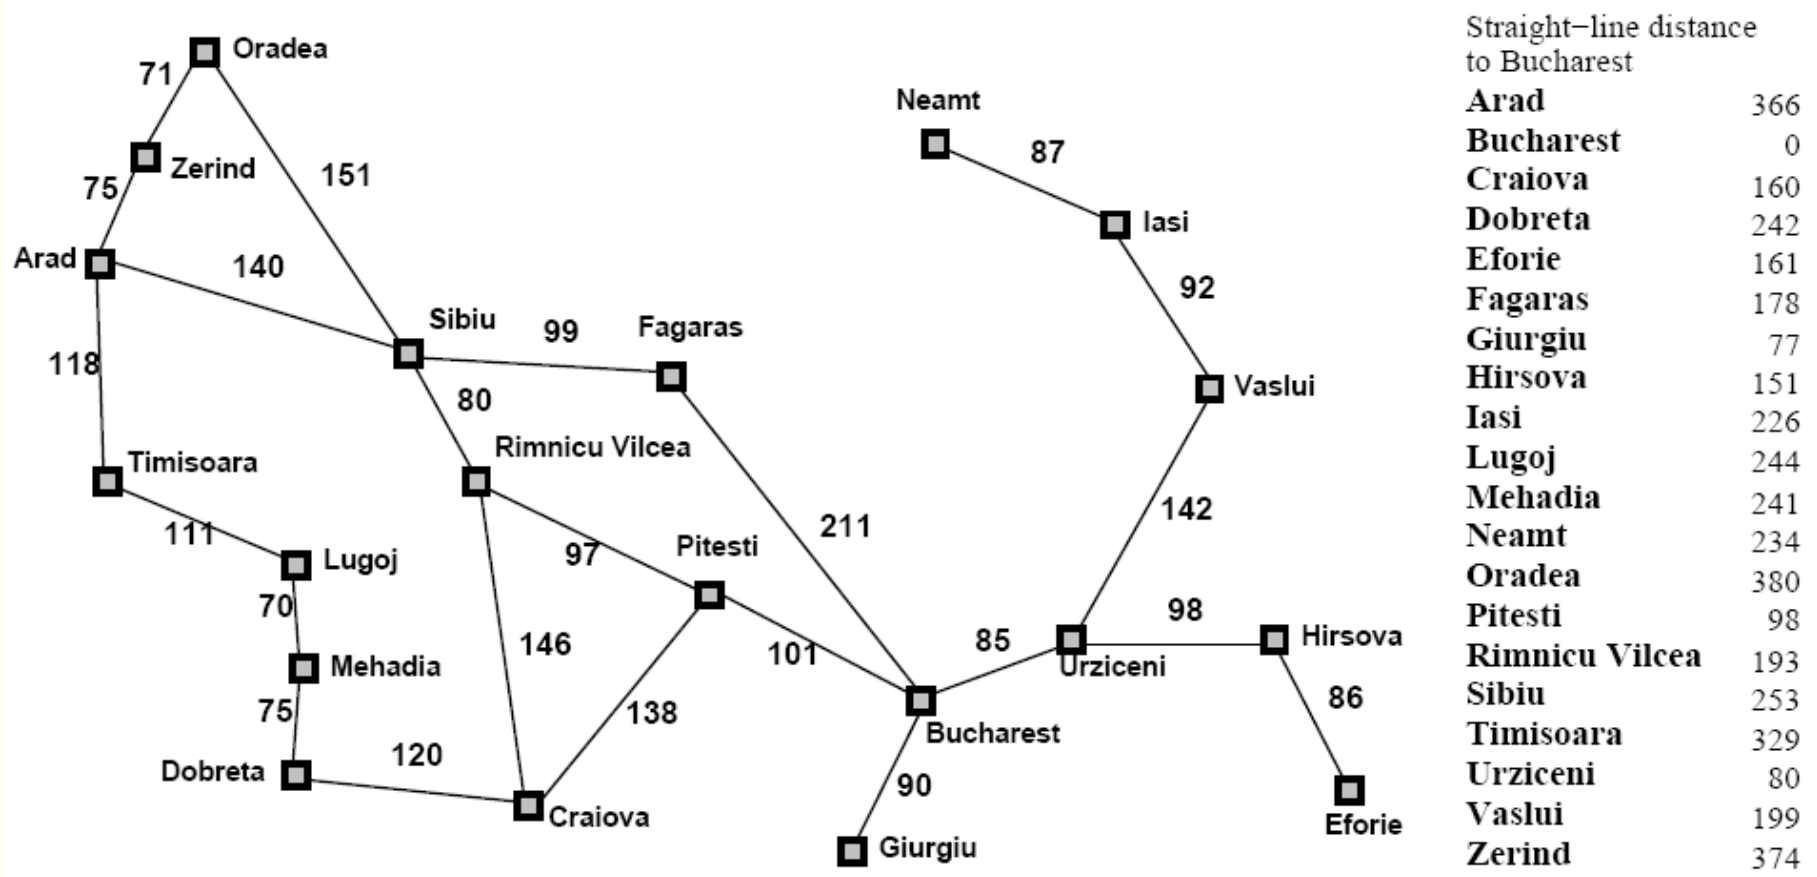
# A\* Search

- When  $h(n) = \text{actual cost to goal}$ 
  - Only nodes in the correct path are expanded
  - Optimal solution is found
- When  $h(n) < \text{actual cost to goal}$ 
  - Additional nodes are expanded
  - Optimal solution is found
- When  $h(n) > \text{actual cost to goal}$ 
  - Optimal solution can be overlooked

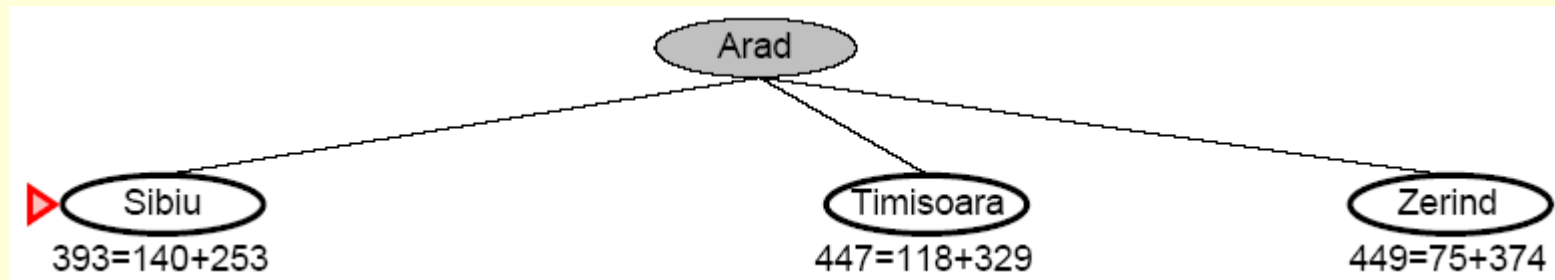

# A\* Search

- A\* is optimal if it uses an **admissible heuristic**
  - $h(n) \leq h^*(n)$  the true cost from node  $n$
  - if  $h(n)$  never overestimates the cost to reach the goal
- Example
  - $h_{SLD}$  never overestimates the actual road distance

# Greedy Best-First Search

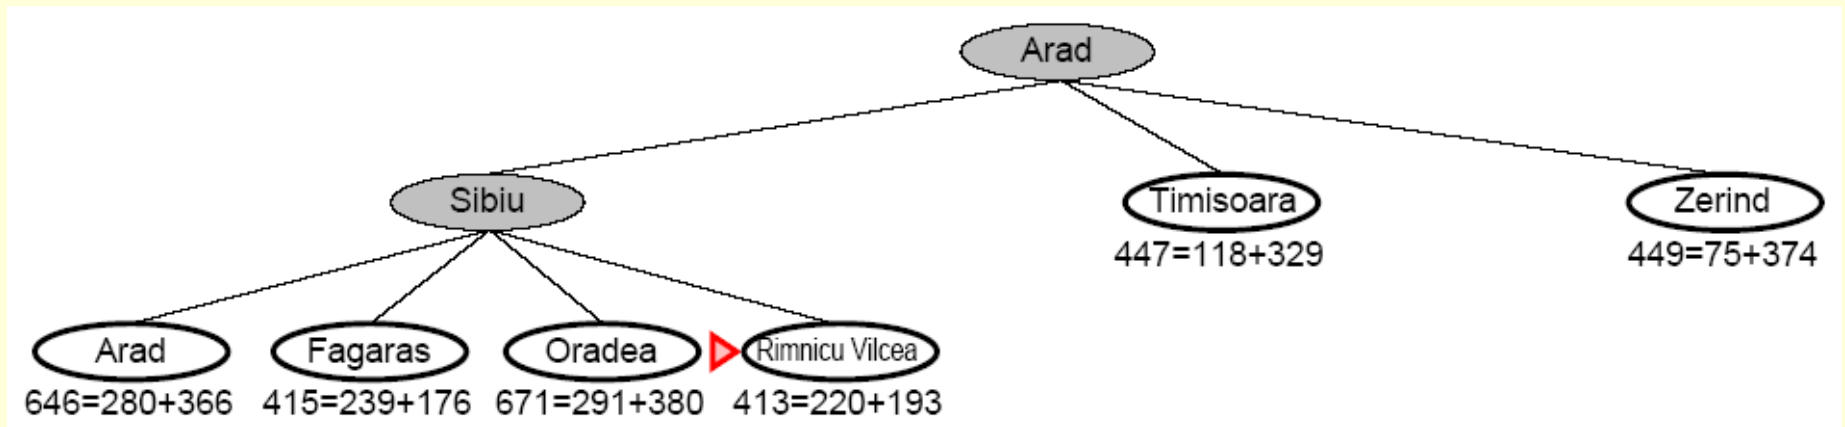


# A\* Search

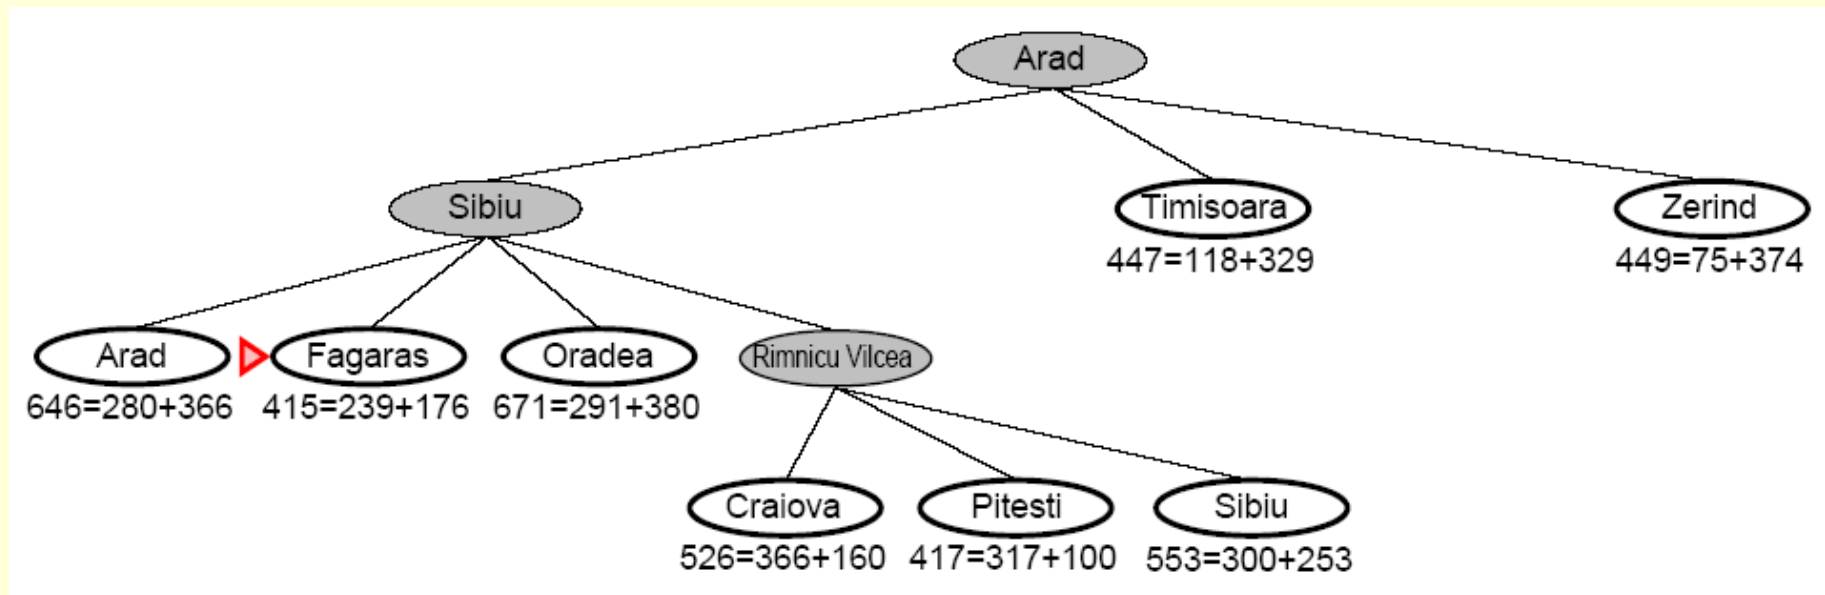




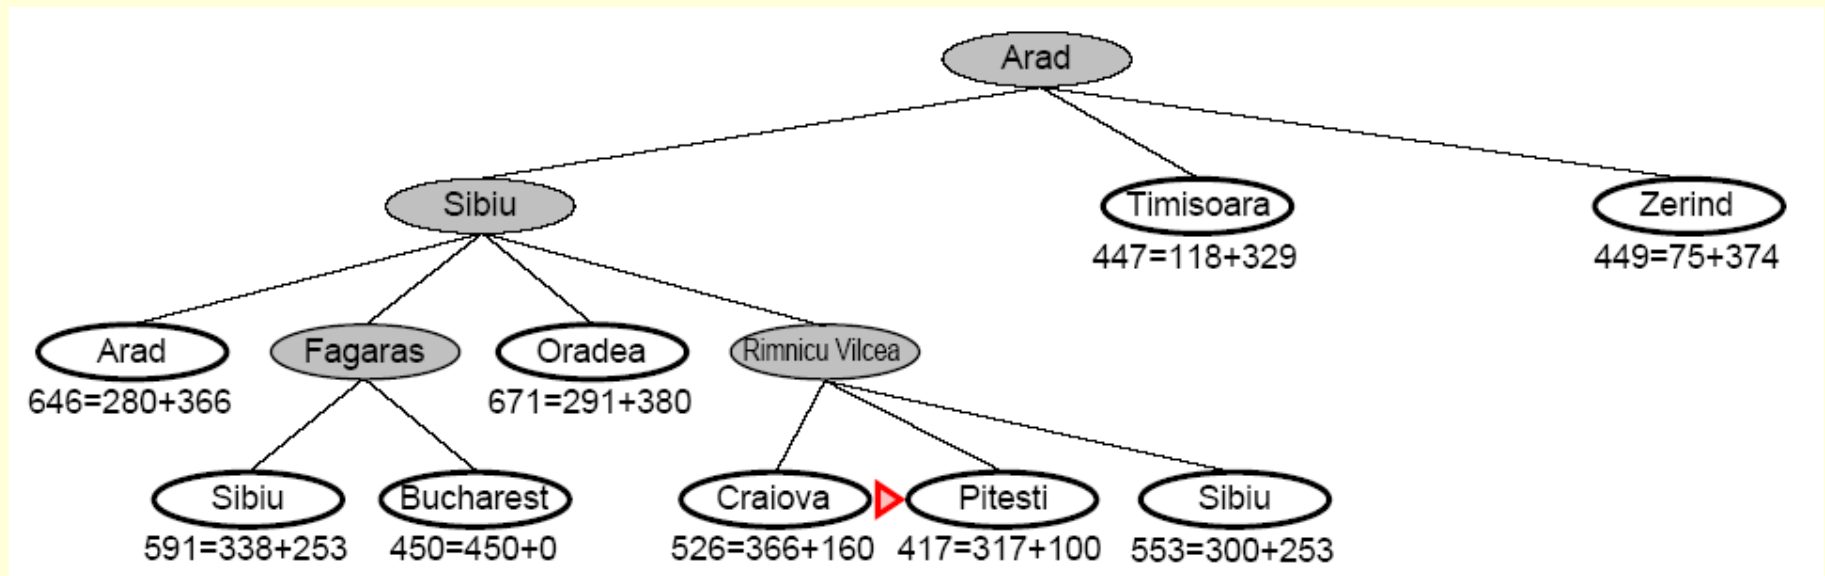
# A\* Search



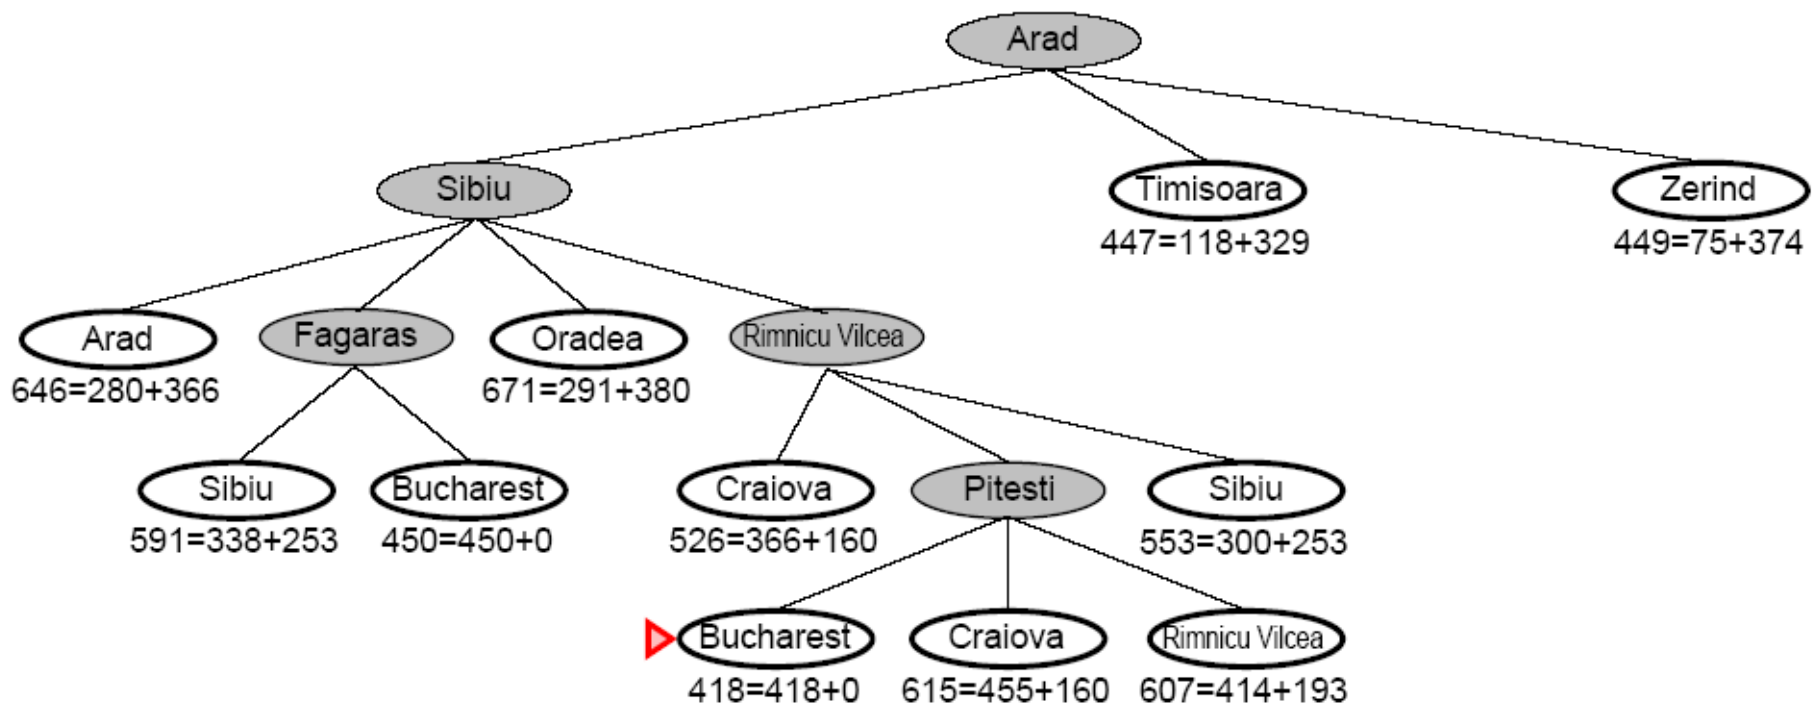
# A\* Search



# A\* Search




# A\* Search



# A\* Search

- A\* expands nodes in increasing  $f$  value
  - Gradually adds  $f$ -contours of nodes (like breadth-first search adding layers)
  - Contour  $f_{i+1}$  where  $f_i < f_{i+1}$



# A\* Search

- Complete
  - Yes, unless there are infinitely many nodes with  $f \leq f(G)$
- Time
  - Exponential in [relative error of  $h$  x length of soln]
  - The better the heuristic, the better the time
    - Best case  $h$  is perfect,  $O(d)$
    - Worst case  $h = 0$ ,  $O(b^d)$  same as BFS
- Space
  - Keeps all nodes in memory and save in case of repetition
  - This is  $O(b^d)$  or worse
  - A\* usually runs out of space before it runs out of time
- Optimal
  - Yes, cannot expand  $f_{i+1}$  unless  $f_i$  is finished

# Memory-Bounded Heuristic Search

- Iterative Deepening A\* (IDA\*)
  - Similar to Iterative Deepening Search, but cut off at  $(g(n) + h(n)) > \text{max}$  instead of  $\text{depth} > \text{max}$
  - At each iteration, cutoff is the first f-cost that exceeds the cost of the node at the previous iteration
- RBFS – see text figures 4.5 and 4.6
- Simple Memory Bounded A\* (SMA\*)
  - Set max to some memory bound
  - If the memory is full, to add a node drop the worst  $(g+h)$  node that is already stored
  - Expands newest best leaf, deletes oldest worst leaf

# Heuristic Functions

- Example: 8-Puzzle
  - Average solution cost for a random puzzle is 22 moves
  - Branching factor is about 3
    - Empty tile in the middle -> four moves
    - Empty tile on the edge -> three moves
    - Empty tile in corner -> two moves
  - $3^{22}$  is approx  $3.1e10$ 
    - Get rid of repeated states
    - 181440 distinct states

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State



# Heuristic Functions

- To use  $A^*$  a heuristic function must be used that never overestimates the number of steps to the goal
- $h1$ =the number of misplaced tiles
- $h2$ =the sum of the Manhattan distances of the tiles from their goal positions

# Heuristic Functions

- $h1 = 7$
- $h2 = 4+0+3+3+1+0+2+1 = 14$

7	2	4
5		6
8	3	1

Start State

1	2	3
4	5	6
7	8	

Goal State

# Dominance

- If  $h_2(n) > h_1(n)$  for all  $n$  (both admissible) then  $h_2(n)$  dominates  $h_1(n)$  and is better for the search
- Take a look at figure 4.8!

# Relaxed Problems

- A Relaxed Problem is a problem with fewer restrictions on the actions
  - The cost of an optimal solution to a relaxed problem is an admissible heuristic for the original problem
- Key point: The optimal solution of a relaxed problem is no greater than the optimal solution of the real problem

# Relaxed Problems

- Example: 8-puzzle
  - Consider only getting tiles 1, 2, 3, and 4 into place
  - If the rules are relaxed such that a tile can move anywhere then  $h1(n)$  gives the shortest solution
  - If the rules are relaxed such that a tile can move to any adjacent square then  $h2(n)$  gives the shortest solution

# Relaxed Problems

- Store sub-problem solutions in a database
  - # patterns is much smaller than the search space
  - Generate database by working backwards from the solution
  - If multiple sub-problems apply, take the max
  - If multiple disjoint sub-problems apply, heuristics can be added

# Learning Heuristics From Experience

- $h(n)$  is an estimate cost of the solution beginning at state  $n$
- How can an agent construct such a function?
- Experience!
  - Have the agent solve many instances of the problem and store the actual cost of  $h(n)$  at some state  $n$
  - Learn from the features of a state that are relevant to the solution, rather than the state itself
    - Generate “many” states with a given feature and determine the average distance
    - Combine the information from multiple features
      - $h(n) = c(1)*x_1(n) + c(2)*x_2(n) + \dots$  where  $x_1, x_2, \dots$  are features

# Optimization Problems

- Instead of considering the whole state space, consider only the current state
- Limits necessary memory; paths not retained
- Amenable to large or continuous (infinite) state spaces where exhaustive search algorithms are not possible
- Local search algorithms can't backtrack



# Local Search Algorithms


- They are useful for solving **optimization problems**
  - Aim is to find a best state according to an **objective function**
- Many optimization problems do not fit the standard search model outlined in chapter 3
  - E.g. There is no goal test or path cost in Darwinian evolution
- State space landscape

# Optimization Problems

- Given measure of goodness (of fit)
  - Find optimal parameters (e.g correspondences)
  - That maximize goodness measure (or minimize badness measure)
- Optimization techniques
  - Direct (closed-form)
  - Search (generate-test)
  - Heuristic search (e.g Hill Climbing)
  - Genetic Algorithm

# Direct Optimization

- The slope of a function at the maximum or minimum is 0
  - Function is neither growing nor shrinking
  - True at global, but also local extreme points
- Find where the slope is zero and you find extrema!
- (If you have the equation, use calculus (first derivative=0))



# Hill Climbing

- Consider all possible successors as “one step” from the current state on the landscape.
- At each iteration, go to
  - The best successor (steepest ascent)
  - Any uphill move (first choice)
  - Any uphill move but steeper is more probable (stochastic)
- All variations get stuck at local maxima

# Hill Climbing


“Like climbing Everest in thick fog with amnesia”

```
function HILL-CLIMBING(problem) returns a state that is a local maximum
  inputs: problem, a problem
  local variables: current, a node
                     neighbor, a node

  current ← MAKE-NODE(INITIAL-STATE[problem])
  loop do
    neighbor ← a highest-valued successor of current
    if VALUE[neighbor] < VALUE[current] then return STATE[current]
    current ← neighbor
  end
```

# Hill Climbing

Problem: depending on initial state, can get stuck on local maxima



In continuous spaces, problems w/ choosing step size, slow convergence

# Hill Climbing

- Local maxima = no uphill step
  - Algorithms on previous slide fail (not complete)
  - Allow “random restart” which is complete, but might take a very long time
- Plateau = all steps equal (flat or shoulder)
  - Must move to equal state to make progress, but no indication of the correct direction
- Ridge = narrow path of maxima, but might have to go down to go up (e.g. diagonal ridge in 4-direction space)

# Simulated Annealing

- Idea: Escape local maxima by allowing some “bad” moves
  - But gradually decreasing their frequency
- Algorithm is randomized:
  - Take a step if random number is less than a value based on both the objective function and the Temperature
- When Temperature is high, chance of going toward a higher value of optimization function  $J(x)$  is greater
- Note higher dimension: “perturb parameter vector” vs. “look at next and previous value”



# Simulated Annealing

```
function SIMULATED-ANNEALING(problem, schedule) returns a solution state
  inputs: problem, a problem
           schedule, a mapping from time to “temperature”
  local variables: current, a node
                    next, a node
                    T, a “temperature” controlling prob. of downward steps

  current  $\leftarrow$  MAKE-NODE(INITIAL-STATE[problem])
  for t  $\leftarrow$  1 to  $\infty$  do
    T  $\leftarrow$  schedule[t]
    if T = 0 then return current
    next  $\leftarrow$  a randomly selected successor of current
     $\Delta E \leftarrow$  VALUE[next] - VALUE[current]
    if  $\Delta E > 0$  then current  $\leftarrow$  next
    else current  $\leftarrow$  next only with probability  $e^{\Delta E/T}$ 
```

# Genetic Algorithms

- Quicker but randomized searching for an optimal parameter vector
- Operations
  - Crossover (2 parents -> 2 children)
  - Mutation (one bit)
- Basic structure
  - Create population
  - Perform crossover & mutation (on fittest)
  - Keep only fittest children

# Genetic Algorithms

- Children carry parts of their parents' data
- Only “good” parents can reproduce
  - Children are at least as “good” as parents?
    - No, but “worse” children don't last long
- Large population allows many “current points” in search
  - Can consider several regions (watersheds) at once

# Genetic Algorithms

- Representation
  - Children (after crossover) should be similar to parent, not random
  - Binary representation of numbers isn't good - what happens when you crossover in the middle of a number?
  - Need "reasonable" breakpoints for crossover (e.g. between R, xcenter and ycenter but not within them)
- "Cover"
  - Population should be large enough to "cover" the range of possibilities
  - Information shouldn't be lost too soon
  - Mutation helps with this issue

# Experimenting With GAs

- Be sure you have a reasonable “goodness” criterion
- Choose a good representation (including methods for crossover and mutation)
- Generate a sufficiently random, large enough population
- Run the algorithm “long enough”
- Find the “winners” among the population
- Variations: multiple populations, keeping vs. not keeping parents, “immigration / emigration”, mutation rate, etc.