

# LECTURE 2: AGENT ARCHITECTURES

**Artificial Intelligence II – Multi-Agent Systems**

**Introduction to Multi-Agent Systems**

**URV, Winter - Spring 2010**

---

# Outline of the talk


- Intelligent agent – reactivity
  - Environments
  - Agent architectures
    - Reactive
    - Deliberative
    - Hybrid
-

# Agent definitions



- "An agent is anything that can be viewed as perceiving its environment through sensors and acting upon that environment through effectors."
- "Autonomous agents are computational systems that inhabit some complex dynamic environment, sense and act autonomously in this environment, and by doing so realize a set of goals or tasks for which they are designed."
- "An autonomous agent is a system situated within and a part of an environment that senses that environment and acts on it, over time, in pursuit of its own agenda and so as to affect what it senses in the future."

# Basic abstract view of an agent



# Reactivity

- An agent has to be able to **react** [adapt its behaviour] in an appropriate way to the dynamic changes in its “**environment**”
  - Other computational agents
  - Human agents/users
  - External information sources (e.g. sensors)
  - Physical objects (e.g. robots)
  - Internet
  - ...
- This is one of several properties that an intelligent agent should have ... **[more on that next week]**

# Kinds of environments (I)



## ■ Accessible vs inaccessible

- ❑ An accessible environment is one in which the agent can obtain complete, accurate, up-to-date information about the environment's state.
- ❑ Most moderately complex environments (including, for example, the everyday physical world and the Internet) are inaccessible.
- ❑ The more accessible an environment is, the simpler it is to build agents to operate in it.

---

# Kinds of environments (II)

## ■ Deterministic vs non-deterministic

- A deterministic environment is one in which any action has a single guaranteed effect — there is no uncertainty about the state that will result from performing an action.
  - The physical world can to all intents and purposes be regarded as non-deterministic.
  - Non-deterministic environments present greater problems for the agent designer.
-

# Kinds of environments (III)

## □ Episodic vs non-episodic

- In an episodic environment, the performance of an agent is dependent on a number of discrete episodes, with no link between the performance of an agent in different scenarios.
- Episodic environments are simpler from the agent developer's perspective because the agent can decide what action to perform based only on the current episode — it need not reason about the interactions between this and future episodes.



# Kinds of environments (IV)

## □ Static vs dynamic

- A static environment is one that can be assumed to remain unchanged except by the performance of actions by the agent.
  - A dynamic environment is one that has other processes operating on it, and which hence changes in ways beyond the agent's control.
  - The physical world is a highly dynamic environment.
-

# Kinds of environments (V)

## □ Discrete vs continuous

- An environment is discrete if there are a fixed, finite number of actions and percepts in it.
- The real world is a continuous environment.

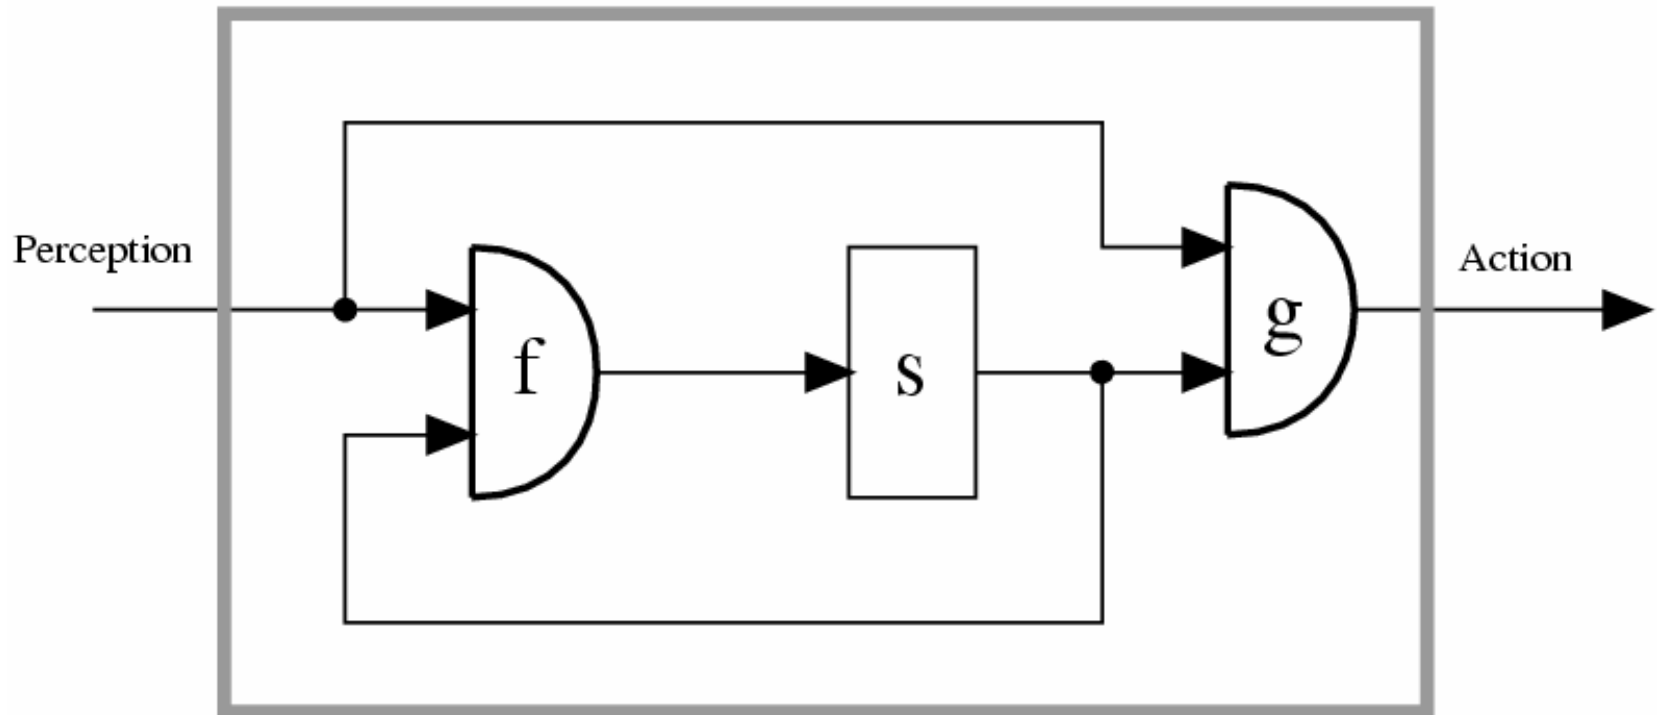
Practical exercise, step 1: think about the kind of environment

---

# Agent architectures

- An **architecture** proposes a particular methodology for building an autonomous agent
  - How the construction of the agent can be decomposed into the construction of a set of **component modules**
  - How these modules should be made to **interact**
  - These two aspects define how the **sensor data** and the **current internal state** of the agent determine the **actions** (effector outputs) and **future internal state** of the agent

# From perception to action



$f$  = state update function

$s$  = internal state

$g$  = output function

---

# Main kinds of agent architectures


- **Reactive** architectures
    - Focused on fast reactions/responses to changes detected in the environment
  - **Deliberative** architectures (symbolic)
    - Focused on long-term planning of actions, centred on a set of basic goals
  - **Hybrid** architectures
    - Combining a reactive side and a deliberative side
-

# Reactive vs Deliberative: example

- Robot that has to reach a certain point
  - Reactive
    - Sensor in the front of the robot
    - Change movement right/left when sensor detects obstacle
      - Minimal computation based on current location and destination point
  - Deliberative
    - Explicit representation of the environment (map)
    - Planning procedure that finds the minimal route between the current position and the destination
      - High computational cost
      - Possible dynamic re-plannings needed

# Reactive Architectures

- There are many unsolved (some would say insoluble) problems associated with **symbolic AI**
  - ❑ Computational cost, brute search
  - ❑ Problems below the 100 ms threshold
    - **For example, face recognition**
- These problems have led some researchers to question the viability of the whole paradigm, and to the development of **reactive** architectures
- Although united by a belief that the assumptions underpinning mainstream AI are in some sense wrong, reactive agent researchers use many different techniques



# Reactive agents – basic ideas

- ❑ Reactive agents have
  - at most a **very simple internal representation** of the world,
  - but provide **tight coupling of perception and action**
- ❑ Behaviour-based paradigm
- ❑ Intelligence is a product of the **interaction** between an agent and its environment



# Classic example: ant colony

- A single ant has very little intelligence, computing power or reasoning abilities
- The union of a set of ants and the interaction between them allows the formation of a highly complex, structured and efficient system.







# SWARM Intelligence

•THE NEXT GENERATION  
OF TECHNOLOGY IS  
MODELED ON INSECTS

by ERIKA D. SMITH  
Boston Journal staff writer

**W**HAT IF computers could make decisions for themselves?

What if you dropped a hatch of computer parts onto a deserted island and they built a civilized communications network all by themselves?

What if your doctor injected a swarm of microchips to attack a hard-to-reach tumor?

A few years ago, these were very big ifs. But a Chagrin Falls startup, named Swarmz Inc., says it is closing in on the technology that could make all of that possible.

Sound scary?

Well, there's no need for a gun-toting Gun Schweenegger just yet. We're not talking artificial intelligence - just swarm intelligence.

"We're taking our inspiration from social insects," said Douglas R. Smith, an Akron resident who co-founded Swarmz.

At the same time, swarm intelligence is a scientific theory

based on the actions of ants, bees and other insects. It asserts complex behavior can emerge from a group of individuals - whether they're bugs or computer nodes - that follow simple rules.

Swarm intelligence isn't exactly new, but it's largely uncharted territory.

Virtually no one has developed a product that applies its principles. But the founders of Swarmz Inc. say they will do it by 2006.

The Chagrin Falls company is writing the mathematical formulas for an all-new, super-efficient wireless network for the military. From there, Smith and company President Mark Hollender hope to parlay their work into other industries, such as manufacturing and supply-chain management.

"They say the possibilities for swarm intelligence products are endless."

"The military has the funding, sense of urgency and need. It's a good place to start," Hollender said. "Later, we can bridge that over to the commercial world."

Biologists long ago noticed the amazing feats social insects are capable of, but swarm

Please see Swarm, D5

# Main characteristics (I)

## ■ Emergent functionality

- Simple agents
- Simple interaction
- Complex behaviour patterns appear as a result of the dynamic interactions
- The global behaviour of the system is not specified a priori
  - Dynamic movement of robots, depending on obstacles

---

# Main characteristics (II)

## ■ Task decomposition

- ❑ Agents composed of autonomous modules
  - ❑ Each module manages a given task
    - Sensor, control, computations
  - ❑ Minimal, low-level communication between modules
  - ❑ There isn't any world global model
  - ❑ There isn't any “planning/controller agent”
-

# Main characteristics (III)

## ■ Raw data

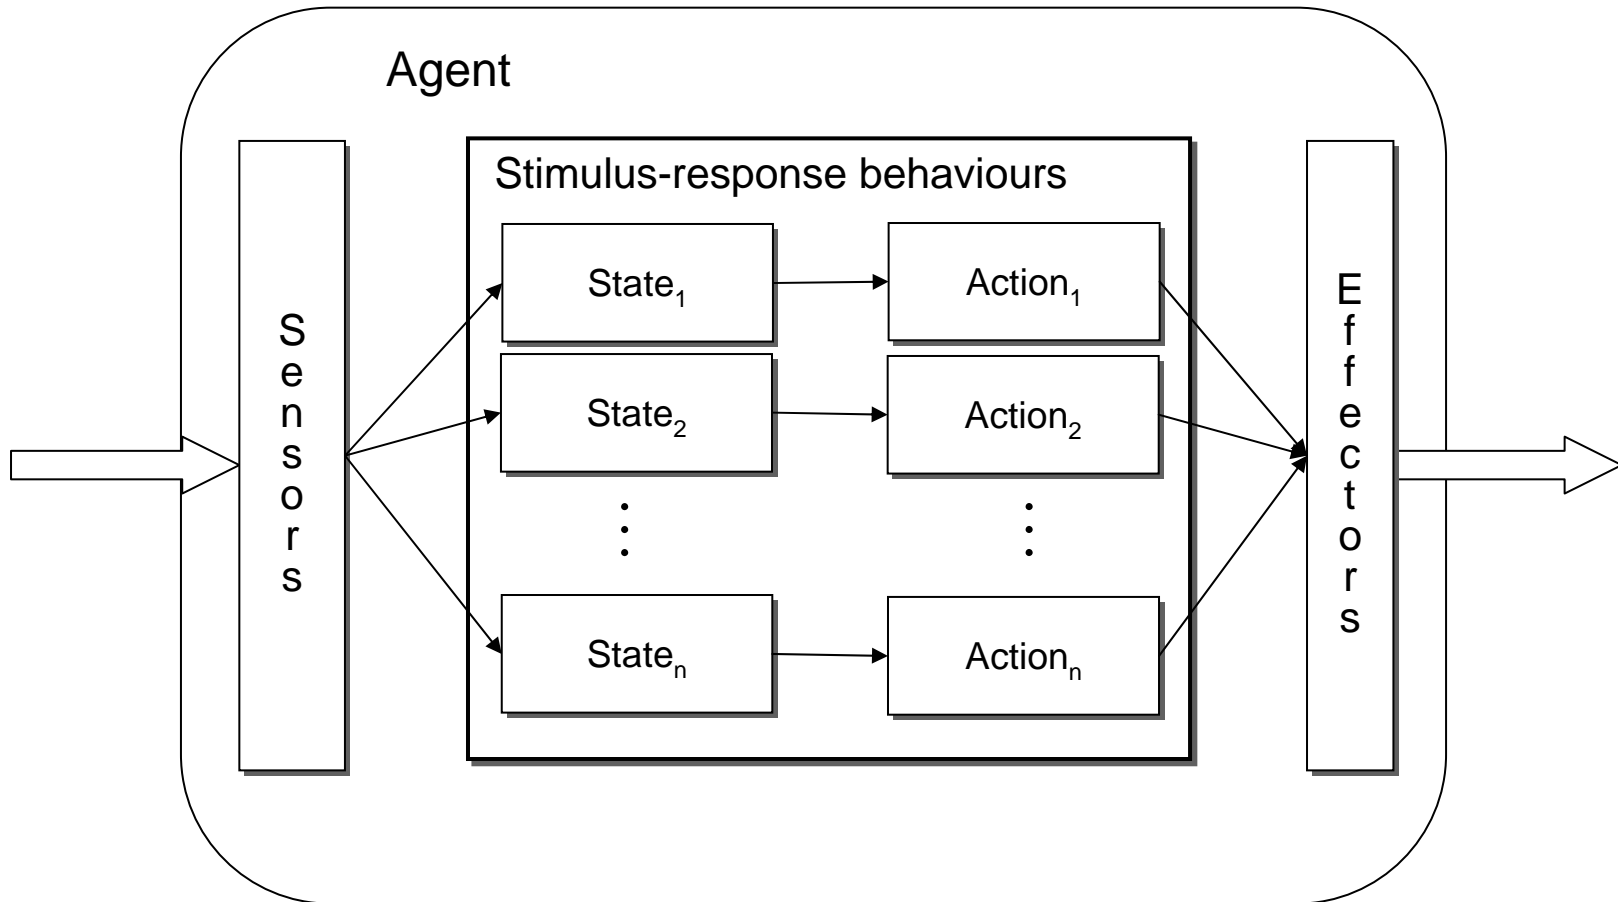
- Basic data from sensors
- There isn't any complex symbolic management of data as in classical AI
  - Refusal of the *Hypothesis of the physic symbols system* [basic pillar of symbolic AI]
    - “Intelligent behaviour can only be obtained in symbol-processing systems”

# Basic concept

- Each behaviour continually maps perceptual input to action output
- Reactive behaviour: **action rules:  $S \rightarrow A$**   
where  $S$  denotes the states of the environment,  
and  $A$  the primitive actions the agent is capable of  
performing.
- Example:

$$\text{action}(s) = \begin{cases} \text{Heater off,} & \text{if temperature is OK in state } s \\ \text{Heater on,} & \text{otherwise} \end{cases}$$

# Basic schema of reactive architecture





---

# Rodney Brooks



Director of the  
Computer  
Science and  
Artificial  
Intelligence Lab  
(MIT)

1997-2007

---

# Brooks refutal of symbolic AI

- Brooks has put forward three theses:
  1. Intelligent behaviour can be generated *without explicit representations* of the kind that symbolic AI proposes
  2. Intelligent behaviour can be generated *without explicit abstract reasoning* of the kind that symbolic AI proposes
    - Reduced computation on sensor-like data
  3. Intelligence is an *emergent* property of certain complex systems

---

# Brooks – key ideas (I)

- **Situatedness:** ‘Real’ intelligence is situated in the world
  - The world is its best model
  - The world is always up-to-date
  - A model is an abstraction, a simplification of the world, considering a particular set of characteristics and disregarding others

# Brooks – key ideas (II)

- ❑ **Embodiment**: ‘Real’ intelligence requires a physical body, and cannot be found in disembodied systems such as theorem provers or expert systems
  - **Physical robots**
- ❑ **Intelligence** and **emergence**: ‘Intelligent’ behavior arises as a result of an agent’s interaction with its environment. Also, intelligence is ‘in the eye of the beholder’; it is not an innate, isolated property

---


# Brooks – behaviour languages

- To illustrate his ideas, Brooks built some systems based on his *subsumption architecture*
  - A subsumption architecture is a hierarchy of task-accomplishing *behaviours*
  - Each behaviour is a rather simple *rule-like structure*
  - Each behaviour ‘competes’ with others to exercise control over the agent, as different behaviours may be applicable at the same time
-

# Behaviour layers

- Lower layers represent more primitive kinds of behaviour (such as avoiding obstacles)
- Higher layers represent more complex behaviours (e.g. identifying an object)
- Lower layers have precedence over layers further up the hierarchy
- The resulting systems are, in terms of the amount of computation they do, *extremely* simple
- Some of the robots do tasks that would be impressive if they were accomplished by symbolic AI systems

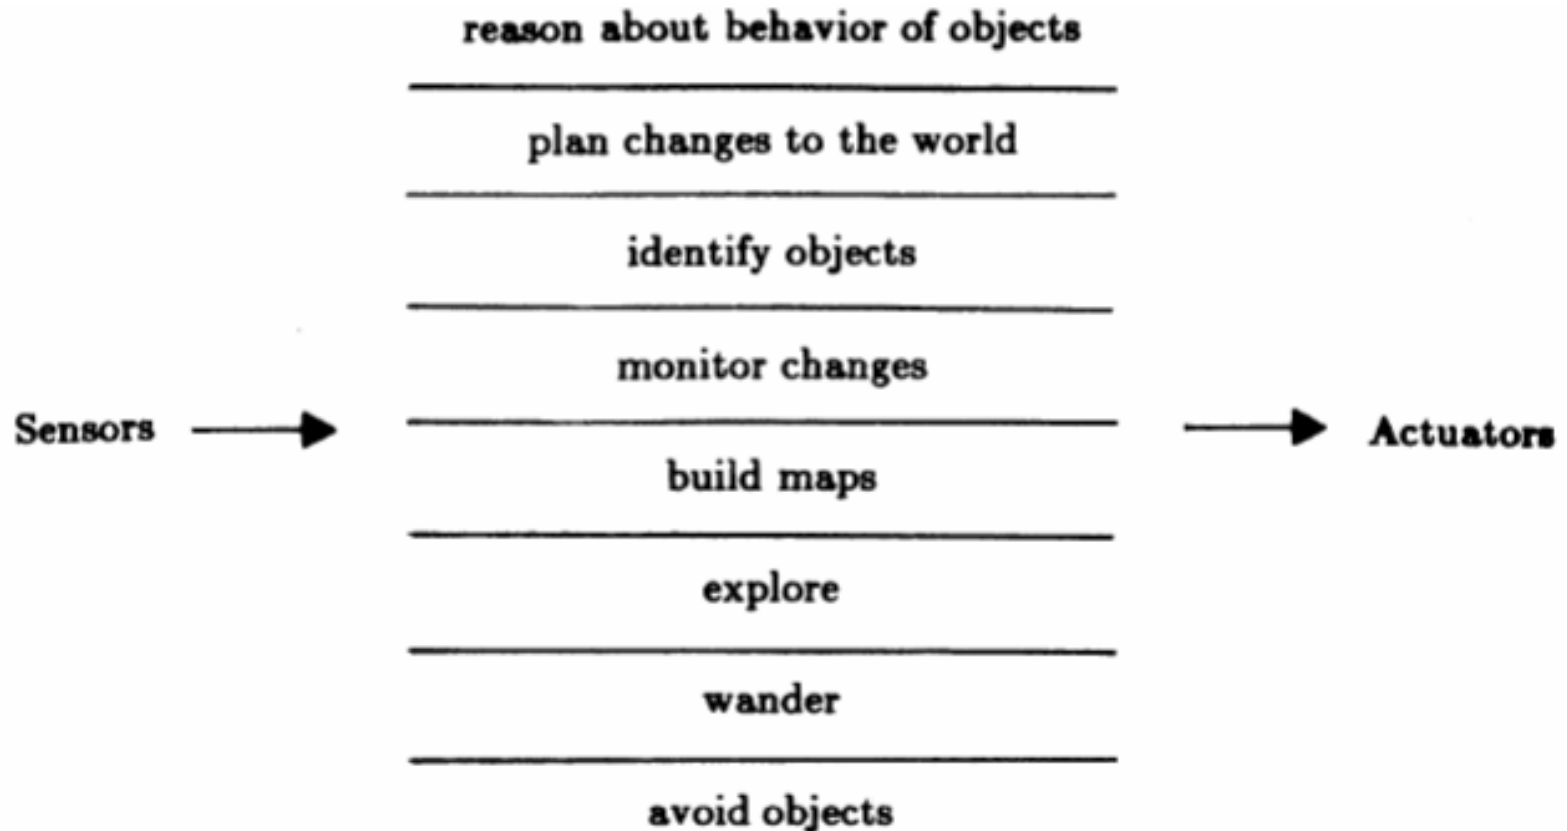
# Decomposition of a Mobile Robot Control System into Functional Modules



Modules in Classical AI systems

From Brooks, "A Robust Layered Control System for a Mobile Robot", 1985


# Decomposition Based on Task Achieving Behaviours



From Brooks, "A Robust Layered Control System for a Mobile Robot", 1985




# Layered Control in the Subsumption Architecture



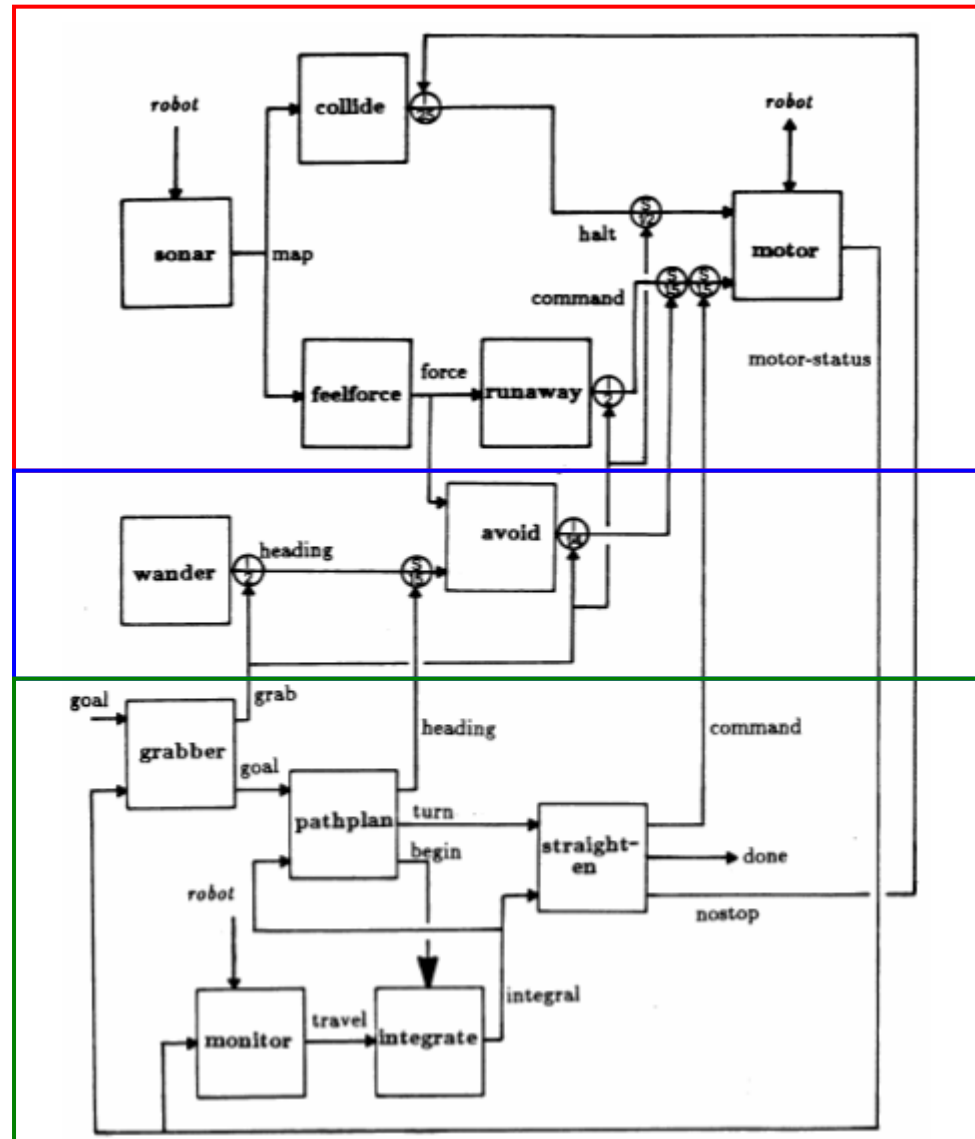
From Brooks, "A Robust Layered Control System for a Mobile Robot", 1985

# Schematic of a Module



From Brooks, "A Robust Layered Control System for a Mobile Robot", 1985

# Levels 0, 1, and 2 Control Systems





Avoid collisions

Random movement

Pick up objects, monitor, explore

From Brooks, "A Robust Layered Control System for a Mobile Robot", 1985



# Steels' Mars Explorer

- Steels' Mars explorer system, using the subsumption architecture, achieves near-optimal cooperative performance in simulated 'rock gathering on Mars' domain:

*The objective is to explore a distant planet, and in particular, to collect sample of a precious rock. The location of the samples is not known in advance, but it **is** known that they tend to be clustered.*



# Steels' Mars Explorer Rules (I)

- For individual (non-cooperative) agents, the lowest-level behaviour, (and hence the behaviour with the highest “priority”) is obstacle avoidance:  
*if detect an obstacle then change direction* (1)
- Any samples carried by agents are dropped back at the mother-ship:  
*if carrying samples and at the base  
then drop samples* (2)
- Agents carrying samples will return to the mother-ship:  
*if carrying samples and not at the base  
then travel up gradient* (3)



---

# Steels' Mars Explorer Rules (II)

- Agents will collect samples they find:  
*if detect a sample then pick sample up* (4)
  - An agent with “nothing better to do” will explore randomly:  
*if true then move randomly* (5)
-

---

# Situated Automata

- A sophisticated theoretical approach is that of Rosenschein and Kaelbling
  - In their *situated automata* paradigm, an agent is specified in a rule-like (declarative) language, and this specification is then compiled down to a *digital machine*, which satisfies the declarative specification
  - This digital machine can operate in a *provable time bound*
  - Reasoning is done *off line*, at *compile time*, rather than *online at run time*
-

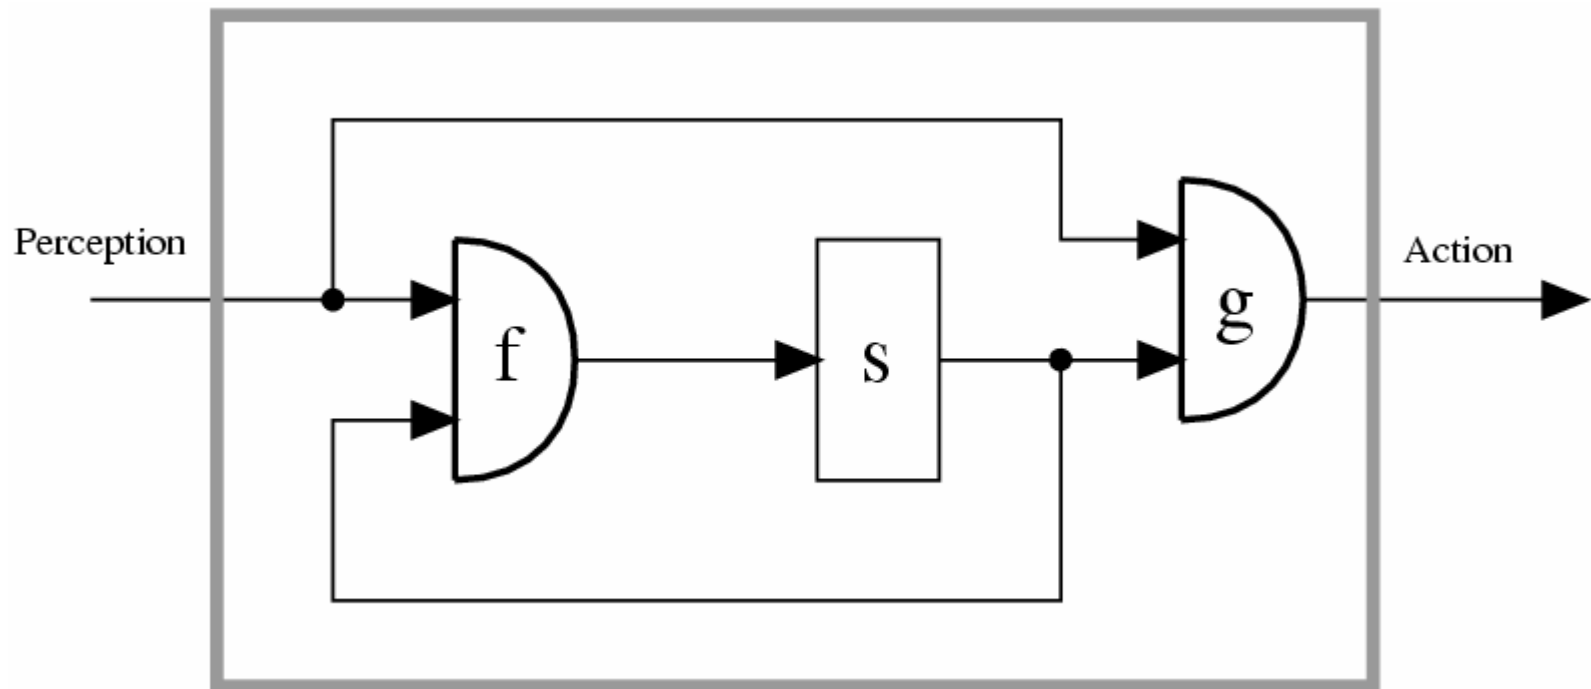


---

# Situated Automata components

- An agent is specified in terms of two components: **perception** and **action**
  - Two programs are then used to synthesize agents
    - RULER is used to specify the perception component of an agent
    - GAPPS is used to specify the action component
-


# Circuit Model of a Finite-State Machine



f = state update function  
s = internal state  
g = output function

From Rosenschein and Kaelbling,  
“A Situated View of Representation and Control”, 1994

# Circuit Model of a Finite-State Machine



From Rosenschein and Kaelbling,  
“A Situated View of Representation and Control”, 1994

# RULER – Situated Automata

- RULER takes as its input three components
  - The **semantics of the agent's inputs** ('whenever bit 1 is on, it is raining')
  - A set of static **facts** ('whenever it is raining, the ground is wet')
  - A specification of the **state transitions** of the world ('if the ground is wet, it stays wet until the sun comes out').
- The programmer then specifies the desired semantics for the output ('if this bit is on, the ground is wet')
- The compiler designs a circuit whose output will have the correct semantics

# GAPPS – Situated Automata

- The GAPPS program takes as its input
  - A set of *goal reduction rules*,
    - Rules that encode information about how goals can be achieved in a given state
  - A top level goal
- Then it generates a program that can be translated into a digital circuit in order to realize the goal
- The generated circuit does not represent or manipulate symbolic expressions; all symbolic manipulation is done at compile time

---

# Advantages of Reactive Agents

- **Simplicity** of individual agents
  - **Flexibility, adaptability**
    - Ideal in very dynamic and unpredictable environments
  - **Computational tractability**
    - Avoiding complex planning/reasoning procedures
    - Avoiding continuous model update
  - **Robustness** against failure
    - No central planning component (e.g. ant colony)
  - **Elegance**
-

# Limitations of Reactive Agents (I)

- Agents without environment models must have sufficient information available from local environment
- If decisions are based on *local* environment, how can we take into account *non-local* information?
  - “Short-term” view
- No long-term planning capabilities
- *Limited applicability*
  - Games, simulations, basic robots (insects)

# Limitations of Reactive Agents (II)

- Difficult to make reactive agents that learn
  - Dynamic evolution of rules?
- Since behaviour emerges from component interactions plus environment, it is hard to see how to *engineer* specific agents (no principled methodology exists)
- It is hard to engineer agents with large numbers of behaviours (dynamics of interactions become too complex to understand)

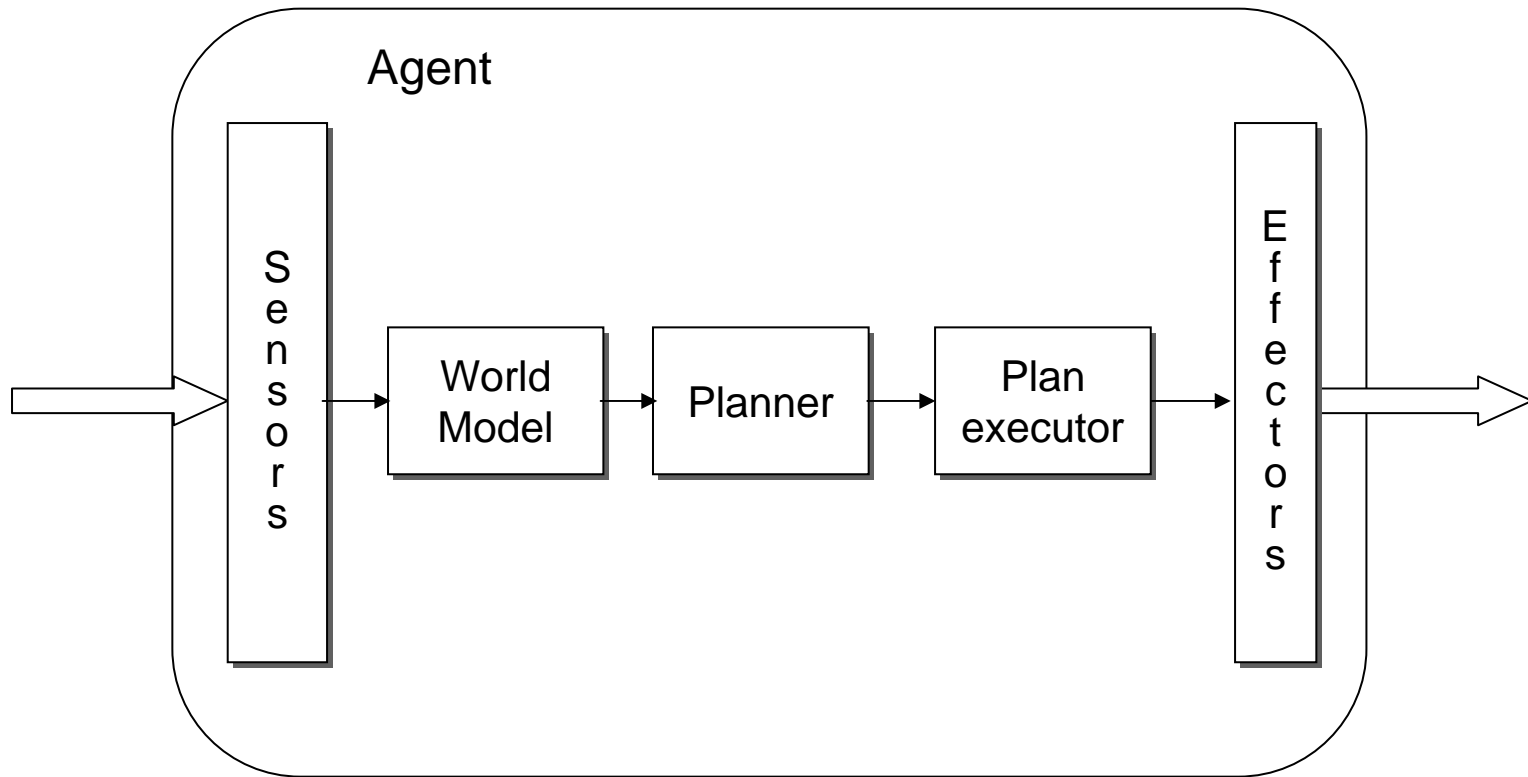


---

# Deliberative agent architecture


- ❑ **Explicit symbolic model** of the world
  - ❑ Decisions are made via logical **reasoning**, based on pattern matching and symbolic manipulation
  - ❑ **Sense-plan-act** problem-solving paradigm of classical AI planning systems
-

# Basic deliberative architecture



# Belief-Desire-Intention (BDI) model

- A theory of practical reasoning.
- Originally developed by Michael E. Bratman in his book "*Intentions, Plans, and Practical Reason*", (1987).
- Concentrates in the roles of the intentions in practical reasoning.



---

# Practical reasoning

- Reasoning directed towards actions — the process of figuring out what to do:

“Practical reasoning is a matter of weighing conflicting **considerations** for and against **competing options**, where the relevant considerations are provided by what the agent **desires** and what the agent **believes**.”  
(Bratman)

“We deliberate not about ends, but about means. We assume the end and consider how and by what **means** it is attained.” (Aristotle)

---

# Human practical reasoning

- Human practical reasoning consists of two activities:
  - *Deliberation*, deciding *what* state of affairs we want to achieve
    - the outputs of deliberation are *intentions*
  - *Means-ends reasoning*, deciding *how* to achieve these states of affairs
    - the outputs of means-ends reasoning are *plans*

# Belief-Desire-Intention paradigm

## ■ Beliefs:

- Agent's view of the environment/world.



## ■ Desires:

- Follow from the beliefs. Desires can be unrealistic and inconsistent.

## ■ Goals:

- A subset of the desires. **Realistic and consistent.** Determine potential processing.

## ■ Intentions:

- A subset of the goals. A goal becomes an intention when an agent decides to commit to it (e.g. by assigning priorities to goals)



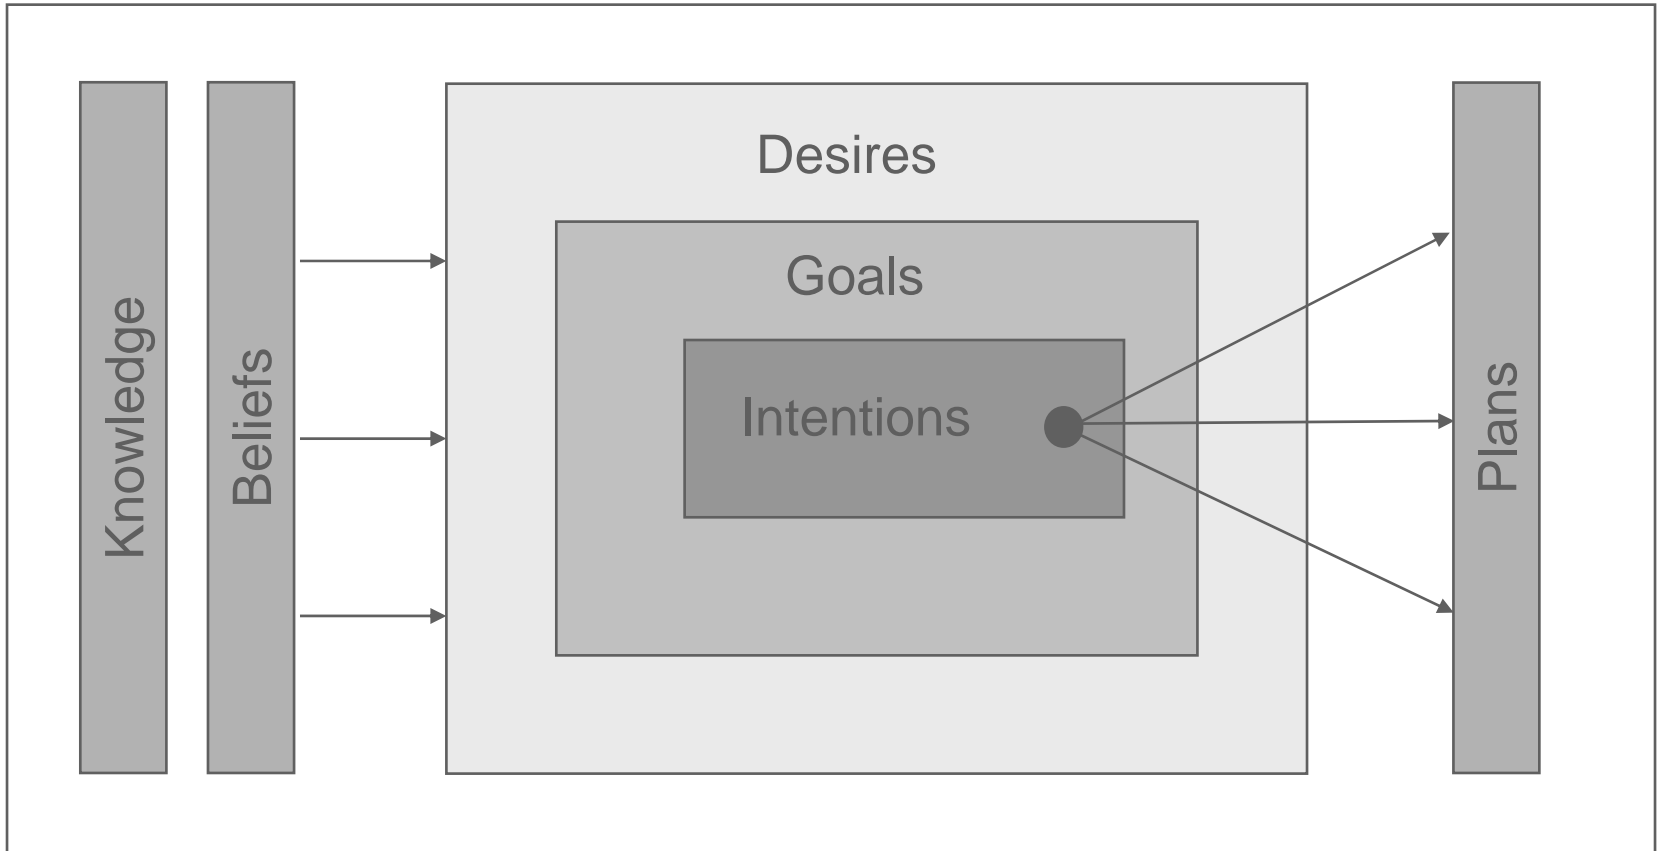
## ■ Plans:

- Sequences of actions that are needed to achieve the intentions, given the agent's beliefs

# BDI plans

- In BDI implementations plans usually have:
  - ❑ a *name*
  - ❑ a *goal*  
invocation condition that is the triggering event for the plan
  - ❑ a *pre-condition list*  
list of facts which must be true for plan to be executed
  - ❑ a *delete list*  
list of facts that are no longer true after plan is performed
  - ❑ an *add list*  
list of facts made true by executing the actions of the plan
  - ❑ a *body*  
list of actions

# Belief-Desire-Intention architecture





# Intention is choice with commitment (Cohen & Levesque)

- There should be "rational balance" among the beliefs, goals, plans, intentions, commitments and actions of autonomous agents.
- Intentions play a big role in maintaining "rational balance"
- An autonomous agent should act on its intentions, not in spite of them
  - ❑ adopt intentions that are feasible
  - ❑ drop the ones that are not feasible
  - ❑ keep (or commit to) intentions, but not forever
  - ❑ discharge those intentions believed to have been satisfied
  - ❑ alter intentions when relevant beliefs change

---

# Using plans to constrain reasoning

- An agent's plans serve to frame its subsequent reasoning problems so as to constrain the amount of resources needed to solve them
    - Agents *commit* to their plans
    - Their plans tell them *what to* reason about, and *what not* to reason about
    - Plans can help reasoning in different levels of abstraction
-

---

# Intention reconsideration

- Intentions (plans) enable the agent to be **goal-driven** rather than event-driven.
  - By committing to intentions the agent can pursue **long-term goals**.
  - However, it is necessary for a BDI agent to reconsider its intentions from time to time:
    - The agent should drop intentions that are no longer achievable.
    - The agent should adopt new intentions that are enabled by opportunities.
-

# Problems in the deliberative approach

## ■ Dynamic world

- Update symbolic world model
- World changes while planning is being done

## ■ Representation language

- Expressive enough to be useful in any domain
- Limited enough to be computationally tractable

## ■ Classical planning => complete, optimal solutions

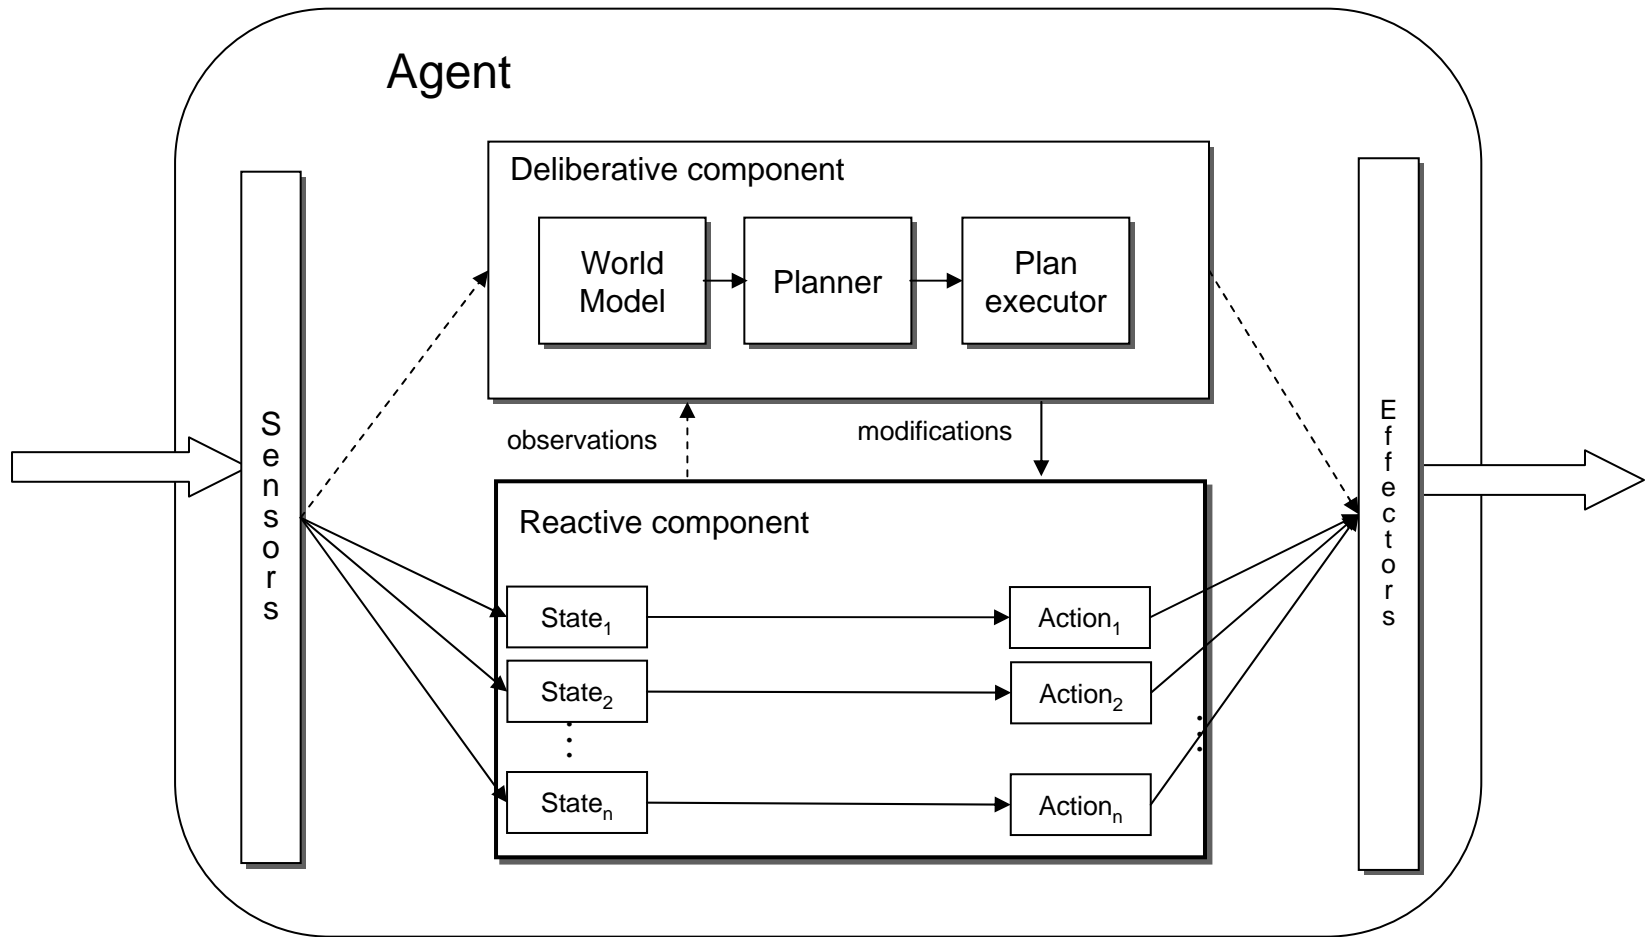
- High computational cost
- Sometimes a sub-optimal low-cost fast reaction can be effective

---

# Hybrid Approaches

- Many researchers have argued that neither a completely deliberative nor a completely reactive approach are suitable for building agents
  - They have suggested using *hybrid* systems, which attempt to marry classical and alternative approaches
  - An obvious approach is to build an agent out of two (or more) subsystems:
    - a *deliberative* one, containing a symbolic world model, which develops plans and makes decisions in the way proposed by symbolic AI
    - a *reactive* one, which is capable of reacting quickly to events without complex reasoning
-

# Hybrid agent architecture



---

# Layered Architectures

- Often, the reactive component is given some kind of precedence over the deliberative one
  - This kind of structuring leads naturally to the idea of a *layered* architecture, of which TOURINGMACHINES and INTERRAP are examples
  - In such an architecture, an agent's control subsystems are arranged into a hierarchy, with higher layers dealing with information at increasing levels of abstraction
-

---


# Layering techniques

- A key problem in such architectures is what kind of control framework to embed the agent's subsystems in, to manage the interactions between the various layers.
  - *Horizontal layering*  
Each layer is directly connected to the sensory input and action output.  
In effect, each layer itself acts like an agent, producing suggestions as to what action to perform.
  - *Vertical layering*  
Sensory input and action output are dealt with by at most one layer each.
-



# Horizontal layering

$m$  possible actions suggested by each layer,  $n$  layers




$O(m^n)$  possible  
options to be  
considered

Introduces bottleneck  
in central control system


(a) Horizontal layering

# Vertical layering

$m$  possible actions suggested by each layer,  $n$  layers



(b) Vertical layering  
(One pass control)



(c) Vertical layering  
(Two pass control)

$O(mn)$  interactions  
between layers

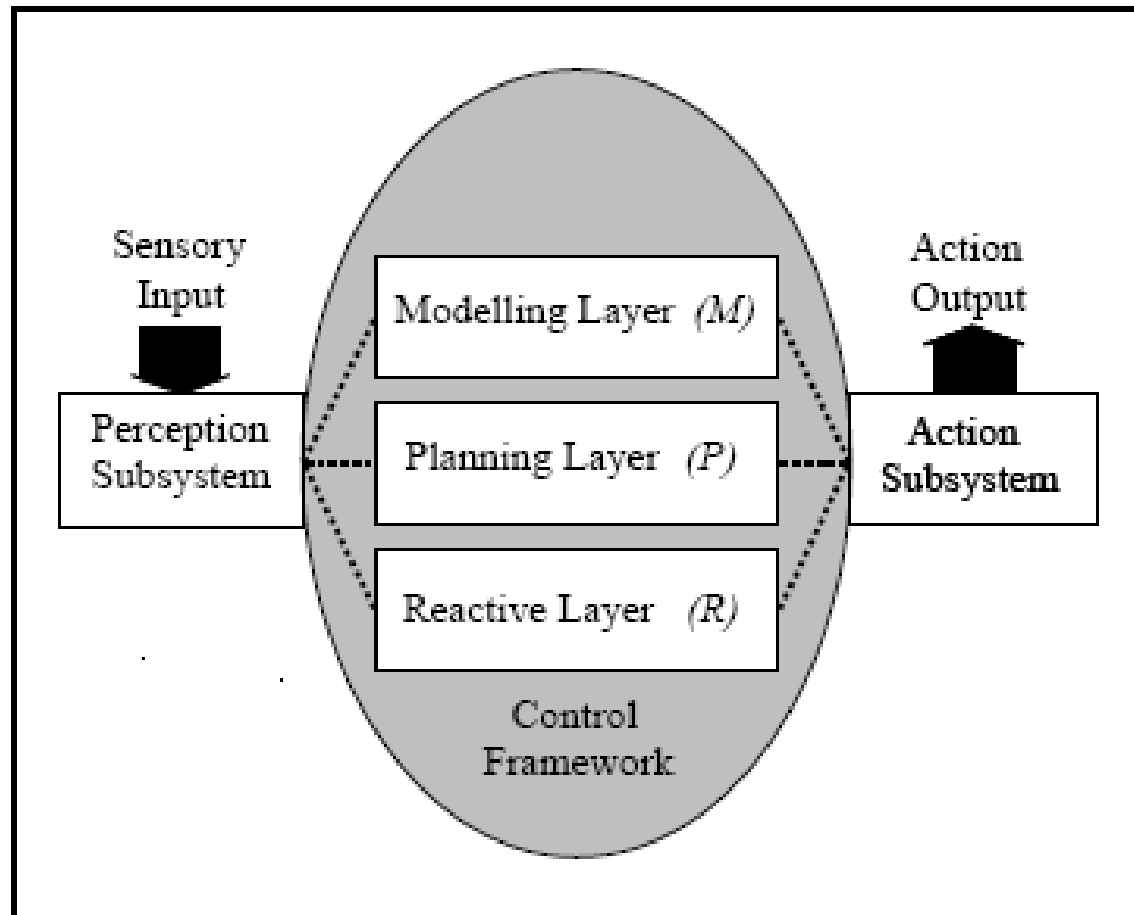
Not fault tolerant to  
layer failure

---

# Ferguson – TOURINGMACHINES

- The TOURINGMACHINES architecture consists of *perception* and *action* subsystems, which interface directly with the agent's environment, and three *control layers*, embedded in a *control framework*, which mediates between the layers
-

# Ferguson – TOURINGMACHINES



# Ferguson – TOURINGMACHINES

- The *reactive layer* is implemented as a set of situation-action rules, *a la* subsumption architecture

rule-1: obstacle-avoidance

if

is-in-front(Obstacle, Observer) and  
speed(Observer) > 0 and  
separation(Obstacle, Observer) < Threshold

then

change-orientation(ObstacleAvoidanceAngle)

- The *planning layer* constructs plans and selects actions to execute in order to achieve the agent's goals

# Ferguson – TOURINGMACHINES

- The *modeling layer* contains symbolic representations of the ‘cognitive state’ of other entities in the agent’s environment
- The three layers communicate with each other and are embedded in a control framework, which use *control rules*

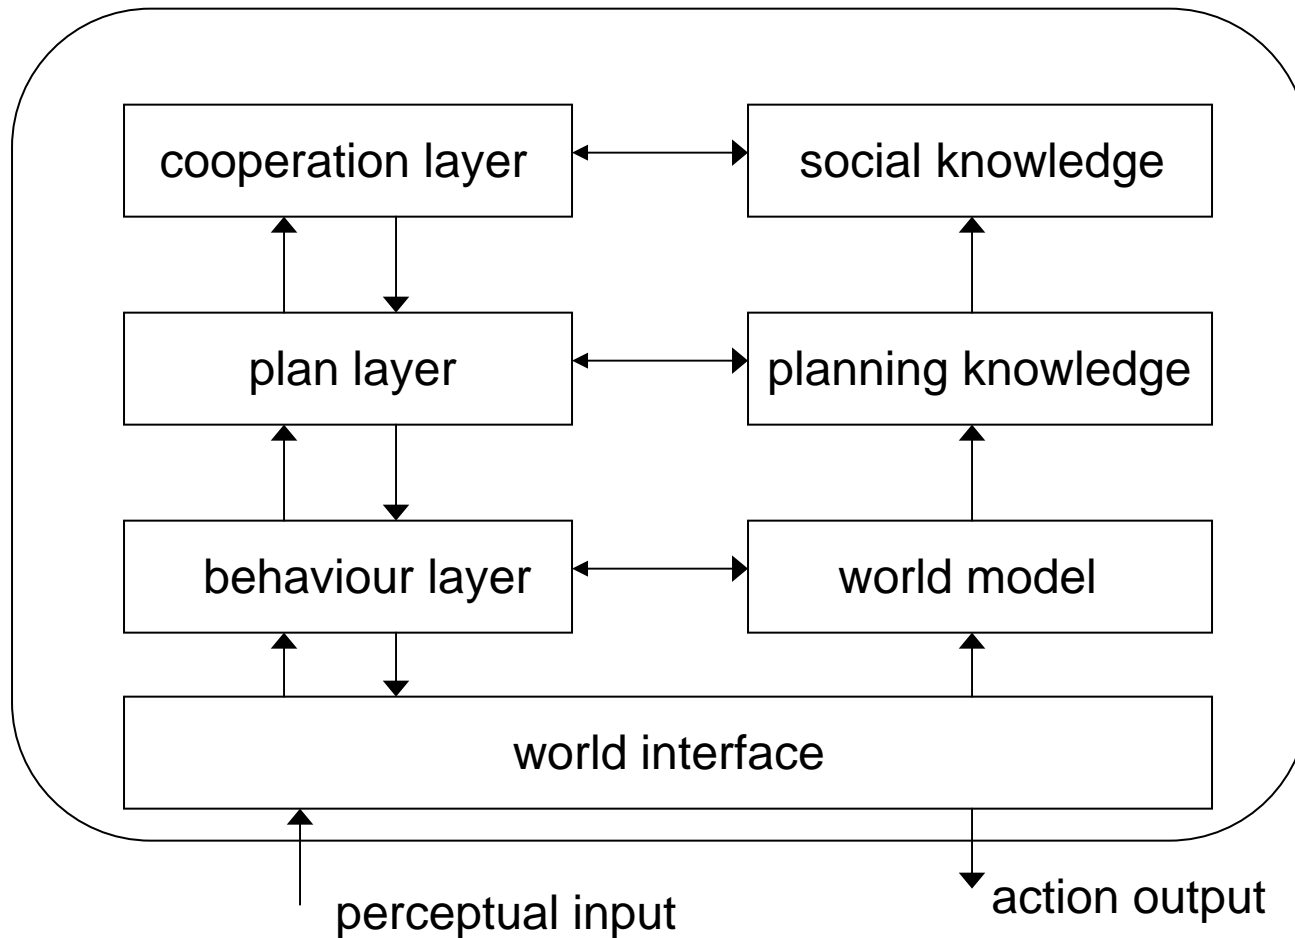
censor-rule-1:

```
if entity(obstacle-6) in perception-buffer  
then remove-sensory-record(layer-R, entity(obstacle-6))
```

[Prevents the reactive layer from reacting in front of obstacle-6]

# Müller –InteRRaP

- Vertically layered, two-pass architecture



# Behaviour layer

- Reactive part of the architecture
- Works with the world model (beliefs on the world state)
- Only level that interacts with the real world
- Has a set of “situation → action” rules
  - Fast recognition of situations that deserve a quick reaction
- Makes routine tasks efficiently, without complex symbolic planning



---

# Planning layer

- Works with the mental model (beliefs on the own agent)
  - Standard deliberative level
  - Implements local behaviour guided towards certain goals
-

---

# Cooperative planning layer

- Works with the social model (beliefs on other agents of the system)
  - Allows planning and cooperation with other agents
    - Global plans of action
    - Conflict resolution
-

---

# Critiques to hybrid architectures

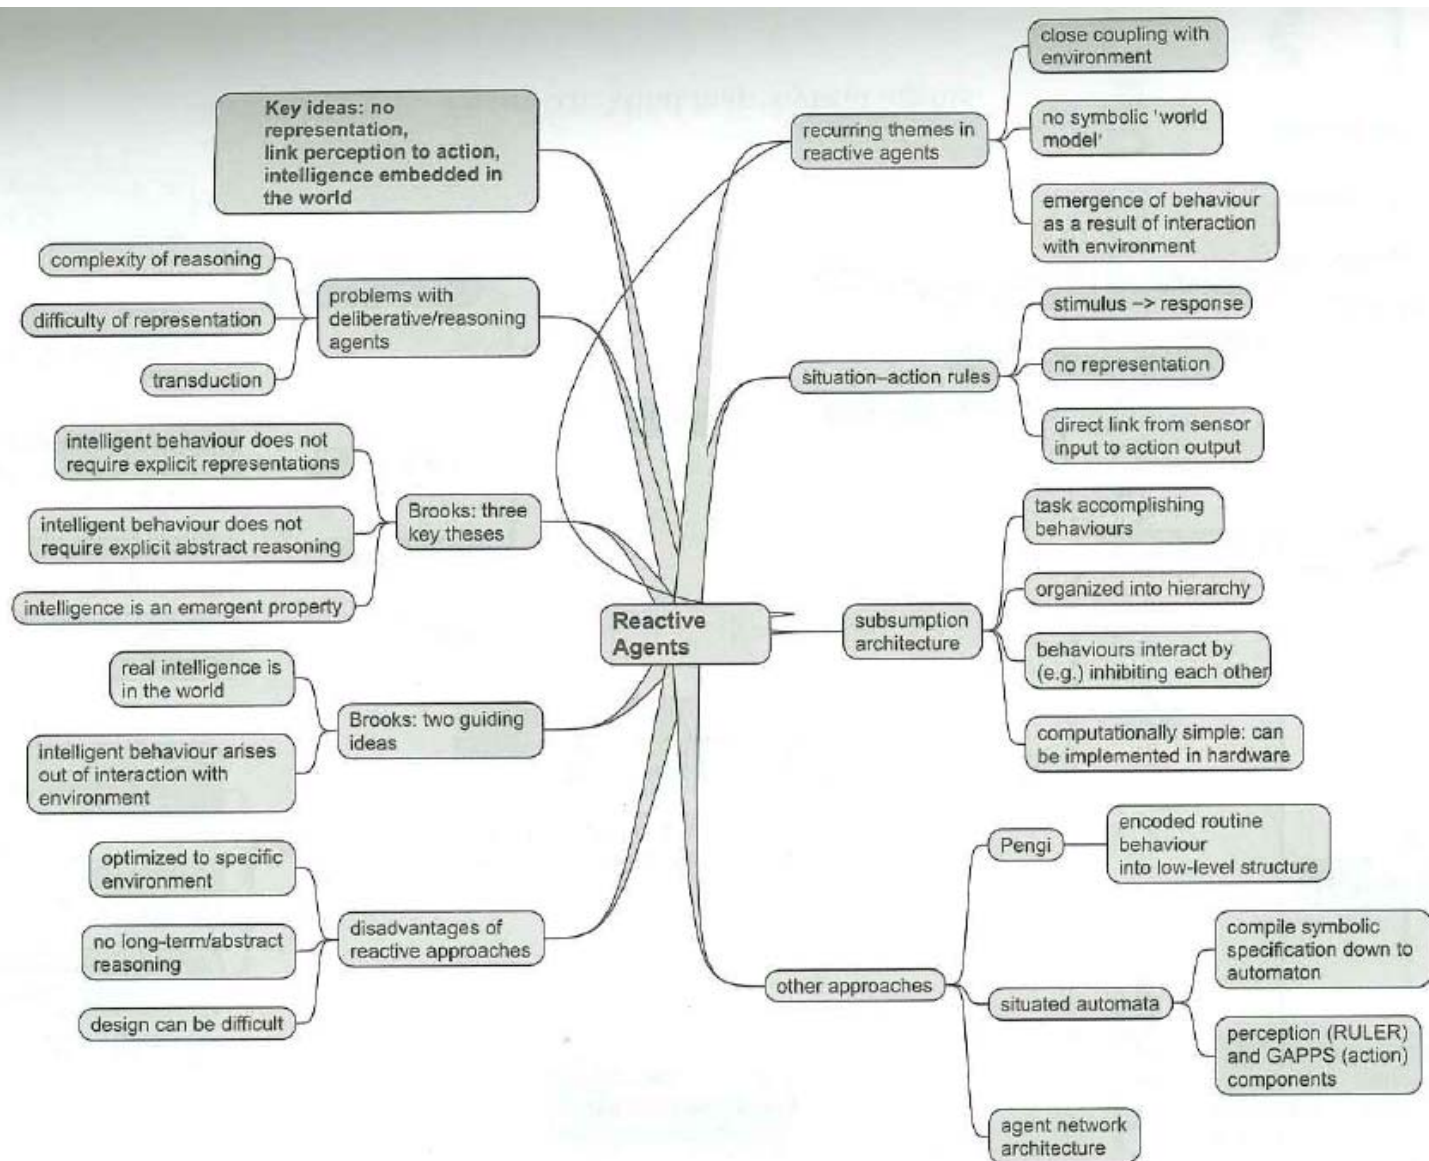
- Lack of general design guiding methodologies
  - Very specific, application dependent
  - Unsupported by formal theories
-

# Readings for this week



- M.Wooldridge: *An introduction to MultiAgent Systems – chapter 5 (reactive, hybrid)*
- A.Mas: *Agentes software y sistemas multi-agente: conceptos, arquitecturas y aplicaciones – chapter 2*
- Articles/book chapters on Moodle web site

# Reactive agents, mind map (Wooldridge)



# Hybrid agents, mind map (Wooldridge)

